

Social Media System Design - Progressive Stages

Stage 2: Search with Elasticsearch

Architecture Overview

Added full-text search capability with Elasticsearch, migrated to NoSQL, and implemented event-driven indexing.

```

graph TB
    subgraph "Stage 2: Search-Enabled Architecture"
        Client[Client/Browser]
        API[API Server<br/>User CRUD<br/>Post CRUD<br/>Search API]

        subgraph "Data Layer"
            MongoDB[(NoSQL Database<br/>MongoDB)]
            ES[(Elasticsearch<br/>Search Index)]
        end

        subgraph "Event Processing"
            Kafka[Apache Kafka<br/>Event Stream]
            Indexer[Indexer Service<br/>Consumes Events]
            Cron[Cron Job<br/>Reindex Service]
        end

        Client -->|HTTP Requests| API
        API -->|Write Operations| MongoDB
        API -->|Search Queries| ES
        API -->|Publish Events| Kafka
        Kafka -->|Subscribe| Indexer
        Indexer -->|Update Index| ES
        MongoDB -->|Read Data| Cron
        Cron -->|Bulk Reindex| ES
    end

    style API fill:#4A90E2
    style MongoDB fill:#47A248
    style ES fill:#FEC514
    style Kafka fill:#231F20
    style Indexer fill:#FF6B6B
    style Cron fill:#9B59B6

```

Data Flow - Create Post with Indexing

```
sequenceDiagram
    participant C as Client
    participant A as API Server
    participant M as MongoDB
    participant K as Kafka
    participant I as Indexer Service
    participant E as Elasticsearch

    Note over C,E: Create Post Flow
```

```
C-->A: POST /api/posts
A-->M: Insert Post Document
M-->A: Success (post_id)
```

```
par Async Indexing
    A-->K: Publish "post.created" Event
    K-->I: Consume Event
    I-->E: Index Post Document
    E-->I: Indexed
end
```

```
A-->C: 201 Created
```

Data Flow - Search Posts

```
sequenceDiagram
    participant C as Client
    participant A as API Server
    participant E as Elasticsearch
    participant M as MongoDB

    Note over C,M: Search Flow
    C->>A: GET /api/posts/search?q=keyword
    A->>E: Query: match(content, keyword)
    E->>E: Rank by relevance,<br/>recency, votes
    E->>A: Post IDs + Scores

    opt Fetch Full Details
        A->>M: Get Posts by IDs
        M->>A: Post Documents
    end

    A->>C: 200 OK + Search Results
```

Reindexing Flow

```
sequenceDiagram
    participant Cron as Cron Job
    participant M as MongoDB
    participant E as Elasticsearch

    Note over Cron,E: Scheduled Reindex (Daily 2 AM)
    Cron->>Cron: Trigger Reindex Job
    Cron->>M: Fetch All Posts (Batch)
    M->>Cron: Posts Data

    loop For Each Batch
        Cron->>E: Bulk Index Documents
        E->>Cron: Batch Indexed
    end

    Cron->>Cron: Log Completion
```

Kafka Topics Structure

```
graph LR
    subgraph "Kafka Topics"
        T1[post.created]
        T2[post.updated]
        T3[post.deleted]
        T4[post.voted]
    end

    API[API Server] -->|Produce| T1
    API -->|Produce| T2
    API -->|Produce| T3
    API -->|Produce| T4

    T1 -->|Consume| Indexer[Indexer Service]
    T2 -->|Consume| Indexer
    T3 -->|Consume| Indexer
    T4 -->|Consume| Indexer

    style T1 fill:#2ECC71
    style T2 fill:#3498DB
    style T3 fill:#E74C3C
    style T4 fill:#F39C12
```

Stage 3: Load Balancing & Caching

Architecture Overview

Production-ready architecture with load balancing, caching layer, and horizontal scaling.

```

graph TB
    subgraph "Stage 3: Scaled Architecture with LB & Cache"
        Client[Clients/Browsers]
        LB[Load Balancer<br/>Nginx/AWS ALB]

        subgraph "Application Layer"
            API1[API Server 1]
            API2[API Server 2]
            API3[API Server 3]
        end

        subgraph "Cache Layer"
            Redis[(Redis Cache<br/>Session & Data)]
        end

        subgraph "Data Layer"
            MongoDB[(MongoDB<br/>Replica Set)]
            ES[(Elasticsearch<br/>Cluster)]
        end

        subgraph "Event Processing"
            Kafka[Kafka Cluster]
            Indexer1[Indexer 1]
            Indexer2[Indexer 2]
            Cron[Cron Service]
        end

        Client -->|HTTPS| LB
        LB -->|Round Robin| API1
        LB -->|Round Robin| API2
        LB -->|Round Robin| API3

        API1 -.→|Check Cache| Redis
        API2 -.→|Check Cache| Redis
        API3 -.→|Check Cache| Redis

        API1 -->|Write/Read| MongoDB
        API2 -->|Write/Read| MongoDB
        API3 -->|Write/Read| MongoDB

        API1 -->|Search| ES
        API2 -->|Search| ES
    
```

```
API3 →|Search| ES
```

```
API1 →|Publish| Kafka
```

```
API2 →|Publish| Kafka
```

```
API3 →|Publish| Kafka
```

```
Kafka →|Consume| Indexer1
```

```
Kafka →|Consume| Indexer2
```

```
Indexer1 →|Update| ES
```

```
Indexer2 →|Update| ES
```

```
MongoDB →|Read| Cron
```

```
Cron →|Reindex| ES
```

```
end
```

```
style LB fill:#E67E22
```

```
style API1 fill:#4A90E2
```

```
style API2 fill:#4A90E2
```

```
style API3 fill:#4A90E2
```

```
style Redis fill:#DC382D
```

```
style MongoDB fill:#47A248
```

```
style ES fill:#FEC514
```

```
style Kafka fill:#231F20
```

Detailed Request Flow with Cache

sequenceDiagram

participant C as Client
participant LB as Load Balancer
participant A as API Server
participant R as Redis Cache
participant M as MongoDB
participant E as Elasticsearch

Note over C,E: Get User Profile (Cache Hit)

C→>LB: GET /api/users/123
LB→>A: Forward Request
A→>R: GET user:123
R→>A: Cache HIT (User Data)
A→>LB: 200 OK + User Data
LB→>C: Response

Note over C,E: Get User Profile (Cache Miss)

C→>LB: GET /api/users/456
LB→>A: Forward Request
A→>R: GET user:456
R→>A: Cache MISS (null)
A→>M: Find user by ID
M→>A: User Document
A→>R: SET user:456 (TTL: 1h)
R→>A: OK
A→>LB: 200 OK + User Data
LB→>C: Response

Note over C,E: Search Posts (with Cache)

C→>LB: GET /api/posts/search?q=tech
LB→>A: Forward Request
A→>R: GET search:tech:page:1
R→>A: Cache MISS
A→>E: Search Query
E→>A: Search Results
A→>R: SET search:tech:page:1 (TTL: 5m)
A→>LB: 200 OK + Results
LB→>C: Response

Cache Strategy Details

```
graph TB
    subgraph "Cache Strategy by Data Type"
        subgraph "User Data"
            U1[User Profile<br/>TTL: 1 hour<br/>Strategy: Cache-Aside]
            U2[User Sessions<br/>TTL: 30 min<br/>Strategy: Write-Through]
        end

        subgraph "Post Data"
            P1[Individual Posts<br/>TTL: 30 min<br/>Strategy: Cache-Aside]
            P2[Post Lists<br/>TTL: 5 min<br/>Strategy: Cache-Aside]
            P3[Trending Posts<br/>TTL: 2 min<br/>Strategy: Refresh-Ahead]
        end

        subgraph "Search Results"
            S1[Search Queries<br/>TTL: 5 min<br/>Strategy: Cache-Aside]
            S2[Popular Searches<br/>TTL: 10 min<br/>Strategy: Cache-Aside]
        end

        subgraph "Computed Data"
            C1[Vote Counts<br/>TTL: 1 min<br/>Strategy: Write-Behind]
            C2[User Stats<br/>TTL: 15 min<br/>Strategy: Cache-Aside]
        end
    end

    style U1 fill:#3498DB
    style U2 fill:#3498DB
    style P1 fill:#2ECC71
    style P2 fill:#2ECC71
    style P3 fill:#2ECC71
    style S1 fill:#F39C12
    style S2 fill:#F39C12
    style C1 fill:#9B59B6
    style C2 fill:#9B59B6
```

Load Balancer Configuration

```
graph LR
    subgraph "Load Balancer Strategies"
        LB[Load Balancer]

        subgraph "Routing Rules"
            R1[/api/users/* → Round Robin]
            R2[/api/posts/* → Round Robin]
            R3[/api/search/* → Least Connections]
            R4[/api/uploads/* → IP Hash]
        end

        subgraph "Health Checks"
            H1[GET /health every 10s]
            H2[Timeout: 5s]
            H3[Unhealthy threshold: 3]
        end

        LB --> R1
        LB --> R2
        LB --> R3
        LB --> R4
        LB -.->|Monitor| H1
        LB -.->|Monitor| H2
        LB -.->|Monitor| H3
    end
```

Write Operation with Cache Invalidation

```
sequenceDiagram
    participant C as Client
    participant LB as Load Balancer
    participant A as API Server
    participant R as Redis
    participant M as MongoDB
    participant K as Kafka
```

Note over C,K: Update Post Flow
C→LB: PUT /api/posts/123
LB→A: Forward Request

```
par Database Update
    A→M: Update Post Document
    M→A: Success
and Cache Invalidation
    A→R: DELETE post:123
    A→R: DELETE post_list:*
    R→A: Keys Deleted
and Event Publishing
    A→K: Publish "post.updated"
    K→A: Acknowledged
end
```

A→LB: 200 OK
LB→C: Response

Note over K: Async: Indexer will update Elasticsearch

Complete Data Flow

```
flowchart TB
    Start([Client Request]) --> LB{Load Balancer}
    LB -->|Route| API[API Server]
    API --> CheckCache{Check Redis Cache}
    CheckCache -->|HIT| ReturnCache[Return Cached Data]
    CheckCache -->|MISS| CheckType{Request Type}
    CheckType -->|Search| ES[Query Elasticsearch]
    CheckType -->|CRUD| DB[Query MongoDB]
    ES --> CacheResult[Cache Search Results]
    DB --> CacheData[Cache DB Results]
    CacheResult --> Return[Return Response]
    CacheData --> Return
    ReturnCache --> Return
    API -->|Write Op| PublishEvent[Publish to Kafka]
    PublishEvent --> Indexer[Indexer Service]
    Indexer --> UpdateES[Update Elasticsearch]
    API -->|Write Op| InvalidateCache[Invalidate Cache]
    Return --> End([Response to Client])

    style LB fill:#E67E22
    style API fill:#4A90E2
    style ES fill:#FEC514
    style DB fill:#47A248
    style CheckCache fill:#DC382D
    style PublishEvent fill:#231F20
```

Comparison: Stage Evolution

Architecture Complexity

```
graph LR
    subgraph "Stage 1"
        S1C[Client] --> S1A[API]
        S1A --> S1D[(DB)]
    end

    subgraph "Stage 2"
        S2C[Client] --> S2A[API]
        S2A --> S2M[(MongoDB)]
        S2A --> S2E[(Elasticsearch)]
        S2A --> S2K[Kafka]
        S2K --> S2I[Indexer]
        S2I --> S2E
    end

    subgraph "Stage 3"
        S3C[Client] --> S3L[LB]
        S3L --> S3A1[API-1]
        S3L --> S3A2[API-2]
        S3A1 --> S3R[(Redis)]
        S3A2 --> S3R
        S3A1 --> S3M[(MongoDB)]
        S3A2 --> S3M
        S3A1 --> S3E[(ES)]
        S3A2 --> S3E
    end

    style S1A fill:#4A90E2
    style S2A fill:#4A90E2
    style S3A1 fill:#4A90E2
    style S3A2 fill:#4A90E2
    style S3L fill:#E67E22
    style S3R fill:#DC382D
```

Performance Metrics Comparison

```
graph TB
    subgraph "Metrics Evolution"
        direction TB

        M1[Stage 1<br/>—<br/>Response Time: 200ms<br/>Throughput: 100 req/s<br/>Availability: 95%]

        M2[Stage 2<br/>—<br/>Response Time: 150ms<br/>Search Time: 50ms<br/>Throughput: 300 req/s<br/>Availability: 97%]

        M3[Stage 3<br/>—<br/>Response Time: 50ms avg<br/>Search Time: 30ms<br/>Throughput: 2000 req/s<br/>Availability: 99.9%]

        M1 -.→|Added Search| M2
        M2 -.→|Added LB & Cache| M3
    end

    style M1 fill:#E74C3C
    style M2 fill:#F39C12
    style M3 fill:#2ECC71
```

Key Features by Stage

FEATURE	STAGE 1	STAGE 2	STAGE 3
User CRUD	✓	✓	✓
Post CRUD	✓	✓	✓
Full-Text Search	✗	✓	✓
Search Ranking	✗	✓ (Relevance, Recency, Votes)	✓
Real-time Indexing	✗	✓ (Kafka)	✓
Batch Reindexing	✗	✓ (Cron)	✓
Load Balancing	✗	✗	✓
Caching	✗	✗	✓ (Redis)
Horizontal Scaling	✗	Partial	✓
High Availability	✗	✗	✓

Technology Stack Summary

Technology Stack			
Layer	Stage 1	Stage 2	Stage 3
Load Bal.	-	-	Nginx/ALB
API	Node.js	Node.js	Node.js (3x)
Cache	-	-	Redis
Database	PostgreSQL	MongoDB	MongoDB Cluster
Search	-	Elasticsearch	ES Cluster
Queue	-	Kafka	Kafka Cluster
Indexer	-	Node.js	Node.js (2x)
Scheduler	-	Cron	Cron Service