

# Caching - 1 Hour Session

---

**Duration:** 60 minutes **Level:** Intermediate



## Session Agenda

- Introduction to Caching (10 min)
- Cache Types & Layers (10 min)
- Caching Strategies & Patterns (15 min)
- Cache Eviction Policies (10 min)
- Implementation & Best Practices (15 min)

## 🎯 Learning Objectives

By the end of this session, you will understand:

- What caching is and when to use it
- Different types and layers of caching
- Common caching patterns and strategies
- Cache eviction policies and their trade-offs
- How to implement and manage caching effectively
- Common pitfalls and how to avoid them

## 1. Introduction to Caching (10 min)

### What is Caching?

**Caching** is the process of storing copies of data in a temporary storage location (cache) for faster access.

 **Quote** “There are only two hard things in Computer Science: cache invalidation and naming things.” - Phil Karlton

## Why Cache?

**Performance Improvements:** -  **Speed:** RAM access ~100ns vs Database ~10ms (100,000x faster) -  **Cost:** Reduce database load and infrastructure costs -  **Scalability:** Handle more requests with same resources -  **User Experience:** Faster response times = happier users

## The Cache Hit/Miss Concept

Request Flow:

1. Check Cache
  - └ HIT: Return cached data (fast!) ✓
  - └ MISS:
    - └ Query database/API (slow)
    - └ Store in cache
    - └ Return data

Cache Hit Ratio = (Cache Hits / Total Requests) × 100%

Target: 80-95% for most applications

## Key Metrics

Metric	Description	Good Target
Hit Ratio	Percentage of requests served from cache	> 80%
Latency	Time to retrieve from cache	< 10ms
TTL	Time data stays in cache	Depends on use case
Memory Usage	Cache size vs available memory	< 80% of allocated

## 2. Cache Types & Layers (10 min)

### Caching Layers (Closest to Farthest from User)

1. Client-Side Cache (Browser)	← Fastest
2. CDN Cache (Edge Locations)	
3. Application Cache (In-Memory)	
4. Distributed Cache (Redis/Memcached)	
5. Database Cache (Query Cache)	
6. Disk Cache than source)	← Slowest (but still faster

### 1. Client-Side Cache

**Browser Cache:** - Stores static assets (images, CSS, JS) - Controlled via HTTP headers

HTTP Headers:

```
Cache-Control: max-age=3600, public
ETag: "33a64df551425fcc55e4d42a148795d9f25f89d4"
Last-Modified: Wed, 21 Oct 2023 07:28:00 GMT
Expires: Thu, 01 Dec 2024 16:00:00 GMT
```

**LocalStorage/SessionStorage:** - JavaScript-accessible storage - 5-10MB limit per domain

```
// LocalStorage example
localStorage.setItem('user_preferences', JSON.stringify(data));
const prefs = JSON.parse(localStorage.getItem('user_preferences'));
```

## 2. CDN Cache

**Content Delivery Network** caches at edge locations worldwide.

```
User in Tokyo → Tokyo Edge Server (cached) ✓ (50ms)  
vs  
User in Tokyo → US Origin Server (not cached) ✗ (250ms)
```

**Use Cases:** - Static assets (images, videos, CSS, JS) - API responses (with proper headers) - Entire pages (for static sites)

## 3. Application-Level Cache

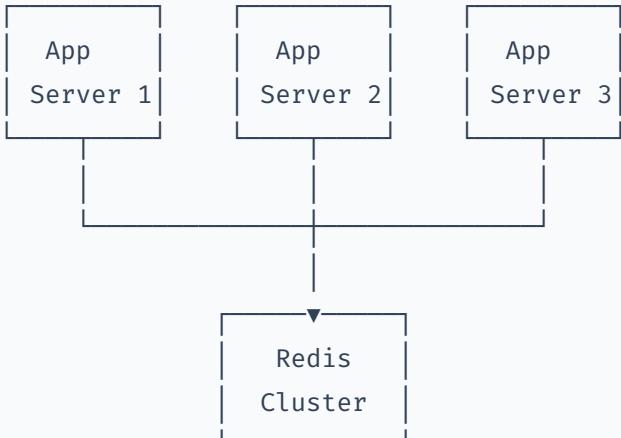
**In-Process Cache:** - Stored in application memory - Fastest but not shared across instances

```
# Python simple in-memory cache  
from functools import lru_cache  
  
@lru_cache(maxsize=1000)  
def get_user(user_id):  
    # This result will be cached  
    return database.query(f"SELECT * FROM users WHERE id={user_id}")
```

## 4. Distributed Cache

**Shared cache across multiple application instances.**

**Popular Solutions:** - **Redis** - Feature-rich, supports data structures - **Memcached** - Simple, fast, key-value only - **Hazelcast** - Distributed in-memory data grid



## 5. Database Cache

**Query Result Cache:** - MySQL Query Cache (deprecated in 8.0) - PostgreSQL shared buffers - MongoDB WiredTiger cache

**ORM-Level Cache:** - Hibernate second-level cache - Django cache framework - Entity Framework caching

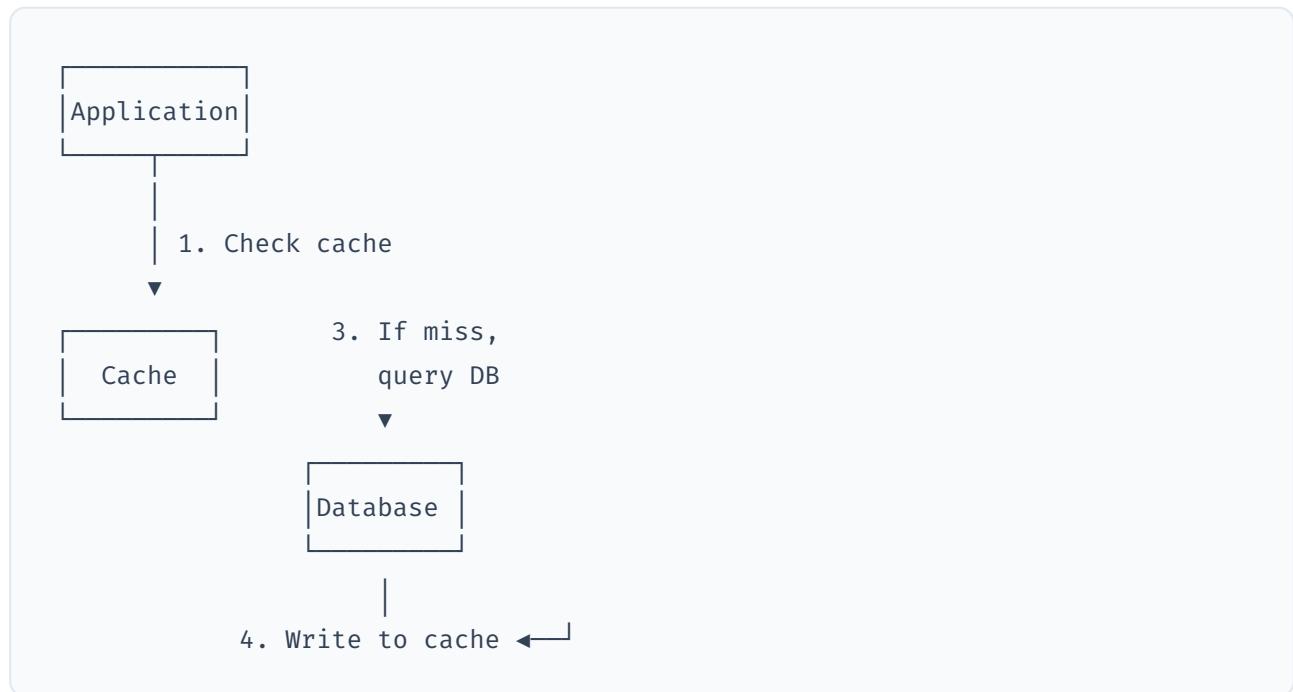
## Cache Type Comparison

Type	Speed	Shared	Complexity	Persistence
Browser	Fastest	No	Low	Session-based
CDN	Very Fast	Yes	Medium	Configurable
In-Memory	Fast	No	Low	No
Redis	Fast	Yes	Medium	Optional
Database	Medium	Yes	Low	Yes

### 3. Caching Strategies & Patterns (15 min)

#### 1. Cache-Aside (Lazy Loading)

**Most common pattern.** Application manages cache explicitly.



#### Implementation:

```
def get_user(user_id):
    # 1. Try cache first
    cache_key = f"user:{user_id}"
    user = cache.get(cache_key)

    if user is not None:
        return user # Cache HIT

    # 2. Cache MISS - query database
    user = database.query(f"SELECT * FROM users WHERE id={user_id}")

    # 3. Store in cache
    cache.set(cache_key, user, ttl=3600) # 1 hour

    return user
```

**Pros:** - Only requested data is cached - Cache failure doesn't break the app

**Cons:** - Cache miss penalty (extra latency) - Data can become stale

## 2. Read-Through Cache

**Cache sits between app and database.** Cache manages data loading.

```
Application → Cache → Database  
(auto-loads on miss)
```

```
# Pseudo-code (handled by cache library)  
user = cache.get(user_id) # Cache automatically queries DB on miss
```

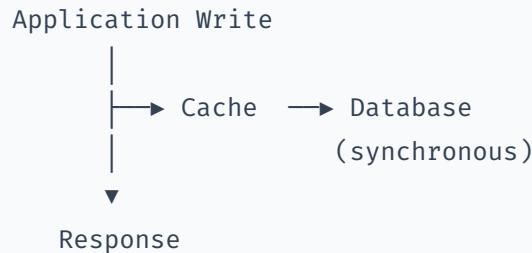
**Pros:** - Simplified application code - Consistent pattern

**Cons:** - Tight coupling to cache - First request still slow

## 3. Write-Through Cache

**Write to cache and database simultaneously.**

```
def update_user(user_id, data):  
    # 1. Update database  
    database.update(user_id, data)  
  
    # 2. Update cache  
    cache_key = f"user:{user_id}"  
    cache.set(cache_key, data, ttl=3600)  
  
    return data
```

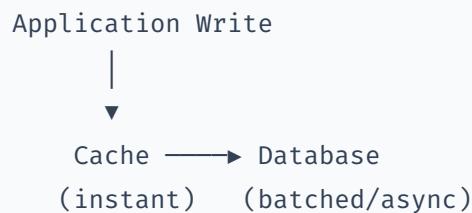


**Pros:** - Cache always consistent with DB - No stale data

**Cons:** - Write latency (two operations) - Wasted writes for unread data

## 4. Write-Behind (Write-Back) Cache

**Write to cache immediately, database asynchronously.**



```

def update_user(user_id, data):
    # 1. Update cache immediately
    cache_key = f"user:{user_id}"
    cache.set(cache_key, data, ttl=3600)

    # 2. Queue database update (async)
    queue.enqueue('db_update', user_id, data)

    return data # Fast response!

```

**Pros:** - Very fast writes - Can batch database operations

**Cons:** - Risk of data loss if cache fails - Complex to implement correctly

## 5. Refresh-Ahead

**Proactively refresh cache before expiration.**

```

import time
from threading import Thread

def get_user_with_refresh(user_id):
    cache_key = f"user:{user_id}"
    user, expiry = cache.get_with_expiry(cache_key)

    if user is None:
        # Cache miss - load from DB
        user = database.query(user_id)
        cache.set(cache_key, user, ttl=3600)
    elif time.time() > (expiry - 300): # 5 min before expiry
        # Refresh in background
        Thread(target=refresh_cache, args=(user_id,)).start()

    return user

def refresh_cache(user_id):
    user = database.query(user_id)
    cache.set(f"user:{user_id}", user, ttl=3600)

```

**Pros:** - Reduced latency (no cache miss delays) - Predictable performance

**Cons:** - Additional complexity - May refresh unused data

## Strategy Comparison

PATTERN	READ PERF	WRITE PERF	COMPLEXITY	DATA FRESHNESS
Cache-Aside	Good	Good	Low	Can be stale
Read-Through	Good	Good	Medium	Can be stale
Write-Through	Good	Slower	Medium	Always fresh
Write-Behind	Good	Fastest	High	Eventually consistent
Refresh-Ahead	Best	Good	High	Very fresh

## 4. Cache Eviction Policies (10 min)

### Why Eviction?

**⚠ Warning** Caches have finite memory. When full, old data must be removed to make room for new data.

### Common Eviction Algorithms

#### 1. LRU (Least Recently Used)

Removes the least recently accessed item.

```
Cache: [A, B, C, D] (max size: 4)
```

```
Access B → [A, C, D, B]
```

```
Access C → [A, D, B, C]
```

```
Add E → [D, B, C, E] (A evicted - oldest access)
```

**Best for:** General purpose, temporal locality

```

from collections import OrderedDict

class LRUCache:
    def __init__(self, capacity):
        self.cache = OrderedDict()
        self.capacity = capacity

    def get(self, key):
        if key not in self.cache:
            return None
        # Move to end (most recently used)
        self.cache.move_to_end(key)
        return self.cache[key]

    def put(self, key, value):
        if key in self.cache:
            self.cache.move_to_end(key)
        self.cache[key] = value
        if len(self.cache) > self.capacity:
            # Remove first item (least recently used)
            self.cache.popitem(last=False)

```

## 2. LFU (Least Frequently Used)

**Removes the least frequently accessed item.**

```

Item A: accessed 10 times
Item B: accessed 5 times
Item C: accessed 2 times ← Evicted first
Item D: accessed 8 times

```

**Best for:** Items with consistent popularity

## 3. FIFO (First In, First Out)

**Removes oldest item regardless of access.**

Cache: [A(1st), B(2nd), C(3rd), D(4th)]

Add E → [B(2nd), C(3rd), D(4th), E(5th)] (A evicted)

**Best for:** Simple use cases, streaming data

## 4. TTL (Time To Live)

**Items expire after a set time.**

```
import time

cache = {}

def set_with_ttl(key, value, ttl_seconds):
    expiry = time.time() + ttl_seconds
    cache[key] = {'value': value, 'expiry': expiry}

def get_with_ttl(key):
    if key not in cache:
        return None

    item = cache[key]
    if time.time() > item['expiry']:
        del cache[key] # Expired
        return None

    return item['value']
```

**Best for:** Time-sensitive data (sessions, tokens)

## 5. Random Replacement

**Randomly selects item to evict.**

**Best for:** When access patterns are unpredictable

## Eviction Policy Comparison

POLICY	COMPLEXITY	HIT RATIO	USE CASE
LRU	Medium	High	General purpose, recency matters
LFU	High	High	Popularity-based content
FIFO	Low	Low	Simple queues
TTL	Low	Medium	Time-sensitive data
Random	Low	Low	Unpredictable patterns

## Redis Eviction Policies

Redis supports multiple policies:

```
# redis.conf
maxmemory 2gb
maxmemory-policy allkeys-lru

# Options:
# noevasion      - Return errors when memory limit reached
# allkeys-lru    - Remove any key, LRU algorithm
# volatile-lru   - Remove keys with TTL, LRU
# allkeys-lfu    - Remove any key, LFU algorithm
# volatile-lfu   - Remove keys with TTL, LFU
# allkeys-random - Remove random key
# volatile-random - Remove random key with TTL
# volatile-ttl   - Remove key with nearest expiration
```

## 5. Implementation & Best Practices (15 min)

### Redis Implementation Examples

#### Basic Operations

```
import redis
import json

# Connect to Redis
r = redis.Redis(host='localhost', port=6379, db=0, decode_responses=True)

# SET with TTL
def cache_user(user_id, user_data):
    key = f"user:{user_id}"
    r.setex(key, 3600, json.dumps(user_data)) # 1 hour TTL

# GET
def get_cached_user(user_id):
    key = f"user:{user_id}"
    data = r.get(key)
    return json.loads(data) if data else None

# DELETE
def invalidate_user_cache(user_id):
    key = f"user:{user_id}"
    r.delete(key)

# CHECK EXISTS
def is_cached(user_id):
    key = f"user:{user_id}"
    return r.exists(key) > 0
```

#### Advanced Patterns

##### 1. Cache Warming:

```

def warm_cache():
    """Pre-load frequently accessed data"""
    popular_users = database.query("SELECT * FROM users ORDER BY views DESC
LIMIT 100")

    pipeline = r.pipeline()
    for user in popular_users:
        key = f"user:{user['id']}"
        pipeline.setex(key, 3600, json.dumps(user))

    pipeline.execute() # Batch operation

```

## 2. Cache Stampede Prevention:

```

import time
import random

def get_with_stampede_protection(key, fetch_func, ttl=3600):
    """Prevent multiple threads from fetching same data simultaneously"""

    data = r.get(key)
    if data:
        return json.loads(data)

    # Try to acquire lock
    lock_key = f"lock:{key}"
    lock_acquired = r.set(lock_key, "1", nx=True, ex=10) # 10 sec lock

    if lock_acquired:
        try:
            # Fetch data
            data = fetch_func()
            r.setex(key, ttl, json.dumps(data))
        return data
    finally:
        r.delete(lock_key)
    else:
        # Wait and retry
        time.sleep(random.uniform(0.1, 0.5))
    return get_with_stampede_protection(key, fetch_func, ttl)

```

### 3. Multi-Level Cache:

```
class MultiLevelCache:
    def __init__(self):
        self.l1_cache = {} # In-memory (fast)
        self.l2_cache = redis.Redis() # Redis (shared)

    def get(self, key):
        # Try L1 first
        if key in self.l1_cache:
            return self.l1_cache[key]

        # Try L2
        data = self.l2_cache.get(key)
        if data:
            # Promote to L1
            self.l1_cache[key] = data
            return data

    return None

def set(self, key, value, ttl=3600):
    # Write to both levels
    self.l1_cache[key] = value
    self.l2_cache.setex(key, ttl, value)
```

## Memcached Implementation

```
import memcache

mc = memcache.Client(['127.0.0.1:11211'], debug=0)

# SET
mc.set("user:123", user_data, time=3600)

# GET
user = mc.get("user:123")

# DELETE
mc.delete("user:123")

# GET MULTIPLE
users = mc.get_multi(["user:123", "user:456", "user:789"])
```

## Application-Level Caching (Node.js)

```
const NodeCache = require('node-cache');
const cache = new NodeCache({ stdTTL: 3600 });

// Middleware for Express
function cacheMiddleware(duration) {
    return (req, res, next) => {
        const key = req.originalUrl;
        const cachedResponse = cache.get(key);

        if (cachedResponse) {
            return res.send(cachedResponse);
        }

        // Override res.send to cache response
        const originalSend = res.send;
        res.send = function(data) {
            cache.set(key, data, duration);
            originalSend.call(this, data);
        };

        next();
    };
}

// Usage
app.get('/api/users', cacheMiddleware(300), (req, res) => {
    // This response will be cached for 5 minutes
    const users = database.getUsers();
    res.json(users);
});
```

# Best Practices

## 1. Cache Key Naming Convention

Pattern: {resource}:{identifier}:{attribute}

Examples:

user:123	- User object
user:123:profile	- User profile
user:123:posts	- User's posts
session:abc123	- Session data
product:456:inventory	- Product inventory
api:v1:users:list:page:1	- API response

## 2. TTL Strategy

```
# Different TTL for different data types
CACHE_TTL = {
    'static_content': 86400,      # 24 hours
    'user_profile': 3600,        # 1 hour
    'user_session': 1800,        # 30 minutes
    'api_response': 300,         # 5 minutes
    'real_time_data': 60,        # 1 minute
}
```

## 3. Cache Invalidation Strategies

**Time-based (TTL):** ~python cache.setex('user:123', 3600, user\_data) # Auto-expires~

**Event-based:** ~python def update\_user(user\_id, new\_data): database.update(user\_id, new\_data) cache.delete(f'user:{user\_id}') # Invalidate on update~

**Tag-based:** ~python # Invalidate all user-related caches tags = cache.smembers(f'tags:user:{user\_id}') for tag in tags: cache.delete(tag)~

## 4. Monitoring & Alerting

```
# Track cache metrics
def cache_get_with_metrics(key):
    start_time = time.time()
    data = cache.get(key)
    latency = time.time() - start_time

    # Log metrics
    if data:
        metrics.increment('cache.hit')
    else:
        metrics.increment('cache.miss')

    metrics.histogram('cache.latency', latency)

    return data
```

**Key Metrics to Monitor:** - Hit/Miss ratio - Cache memory usage - Eviction rate - Average latency - Connection pool usage

## 5. Cache Size Calculation

```
import sys

def calculate_cache_size(data_size_bytes, num_items, overhead=1.2):
    """
    data_size_bytes: Average size of each cached item
    num_items: Number of items to cache
    overhead: Redis overhead factor (keys, metadata)
    """
    total_size = data_size_bytes * num_items * overhead

    # Convert to GB
    size_gb = total_size / (1024 ** 3)

    return size_gb

# Example: Cache 1M user profiles, avg 2KB each
cache_size = calculate_cache_size(2048, 1_000_000)
print(f"Required cache size: {cache_size:.2f} GB")
```

## 💡 Common Pitfalls & Solutions

### 1. Cache Stampede

**Problem:** Multiple requests for expired key hit database simultaneously.

Cache expires → 1000 concurrent requests → 1000 DB queries! 💥

**Solution:** Use locks or probabilistic early expiration

```

def get_with_early_expiration(key, ttl=3600, beta=1.0):
    data, expiry, delta = cache.get_with_metadata(key)

    if data is None:
        # Cache miss - fetch and store
        data = fetch_from_db(key)
        cache.set(key, data, ttl)
        return data

    # Probabilistic early refresh
    time_to_expiry = expiry - time.time()
    if time_to_expiry < 0:
        return data

    # XFetch algorithm
    if time.time() - delta * beta * math.log(random.random()) ≥ expiry:
        # Refresh in background
        refresh_async(key)

    return data

```

## 2. Stale Data

**Problem:** Cache returns outdated information.

**Solutions:** - Shorter TTL for frequently changing data - Event-driven invalidation - Write-through caching - Cache versioning

```

# Cache versioning
CACHE_VERSION = "v2"

def versioned_key(key):
    return f"{CACHE_VERSION}:{key}"

# When data structure changes, increment version
# Old cached data automatically becomes inaccessible

```

## 3. Cache Penetration

**Problem:** Requests for non-existent keys always hit database.

```
Request for user:99999 (doesn't exist)
→ Cache miss
→ Database query (returns null)
→ No caching (null not cached)
→ Repeat 1000 times 💥
```

**Solution:** Cache null results with short TTL

```
def get_user(user_id):
    key = f"user:{user_id}"
    user = cache.get(key)

    if user == "NULL": # Cached null result
        return None

    if user:
        return user

    user = database.query(user_id)

    if user is None:
        cache.setex(key, 60, "NULL") # Cache null for 1 min
    else:
        cache.setex(key, 3600, user)

    return user
```

## 4. Memory Overflow

**Problem:** Cache grows too large, OOM errors.

**Solutions:** - Set maxmemory limit - Configure eviction policy - Monitor memory usage - Use TTL aggressively

```
# redis.conf
maxmemory 2gb
maxmemory-policy allkeys-lru
```

## 5. Cache Inconsistency

**Problem:** Cache and database out of sync.

**Solutions:**

```
# Use database transactions with cache updates
def update_user_transactional(user_id, data):
    try:
        # Start transaction
        db.begin()
        db.update('users', user_id, data)
        db.commit()


        # Only update cache after successful DB update
        cache.set(f'user:{user_id}', data, ttl=3600)

    except Exception as e:
        db.rollback()
        # Don't update cache if DB update failed
        raise e
```

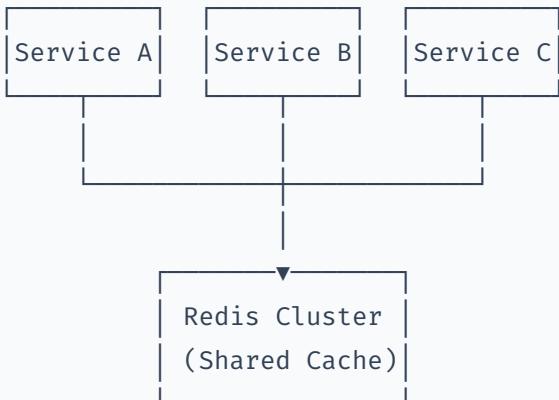


## Cache Architecture Patterns

### Pattern 1: Simple Web Application



### Pattern 2: Microservices with Shared Cache



## Pattern 3: Multi-Tier Caching

Browser Cache (L1)

↓

CDN Cache (L2)

↓

App Cache (L3)

↓

Redis Cache (L4)

↓

DB Query Cache (L5)


↓

Database



## Decision Framework

### When to Cache?



**Cache if:** - ✓ Read frequency >> Write frequency - ✓ Expensive to compute/fetch - ✓ Can tolerate some staleness - ✓ Limited data size

**Don't cache if:** - ✗ Write-heavy workload - ✗ Must be real-time accurate - ✗ Data is already fast to fetch - ✗ Unlimited data size

## Choosing Cache Type

REQUIREMENT	RECOMMENDED SOLUTION
Static assets	CDN (CloudFront, Cloudflare)
Session data	Redis with TTL
Database query results	Application cache + Redis
API responses	CDN + Redis
Computed values	Redis with refresh-ahead
User preferences	Browser localStorage + Redis
Real-time data	Redis with short TTL (30-60s)



## Key Takeaways

1. **Caching dramatically improves performance** but adds complexity
2. **Choose the right cache layer** based on data access patterns
3. **Cache invalidation is hard** - use TTL, event-based, or hybrid approaches
4. **Monitor your cache** - hit ratio, memory, latency
5. **Handle edge cases** - stampede, penetration, inconsistency
6. **Different data needs different strategies** - one size doesn't fit all
7. **Start simple** - add complexity only when needed



## Practical Exercise

**Scenario:** You're building a social media feed application.

**Requirements:** - User posts (10k posts/day) - User profiles (1M users) - Feed generation (complex algorithm, 2s to compute) - Real-time likes/comments

**Questions:** 1. What should you cache? 2. What TTL for each cache? 3. Which caching strategy for each? 4. How to handle real-time updates?

**Suggested Solution:**

**Cache Strategy:**

**User Profiles:**

- **Strategy:** Cache-Aside
- **TTL:** 1 hour
- **Invalidation:** On profile update

**Feed (Personalized):**

- **Strategy:** Cache-Aside with refresh-ahead
- **TTL:** 5 minutes
- **Invalidation:** Time-based + new post event

**Feed (Public):**

- **Strategy:** Write-through
- **TTL:** 1 minute
- **Invalidation:** On new post

**Posts:**

- **Strategy:** Write-through
- **TTL:** 1 hour
- **Invalidation:** On edit/delete

**Likes/Comments Count:**

- **Strategy:** Write-behind (async update)
- **TTL:** 30 seconds
- **Invalidation:** Eventual consistency OK



## Redis Commands Cheatsheet

```
# String operations
SET key value EX 3600          # Set with TTL
GET key                         # Get value
DEL key                          # Delete
EXISTS key                       # Check exists
TTL key                          # Check remaining TTL
EXPIRE key 3600                 # Set TTL on existing key

# Hash operations (for objects)
HSET user:123 name "John"
HGET user:123 name
HGETALL user:123
HDEL user:123 name

# List operations
LPUSH list value                # Add to front
RPUSH list value                # Add to end
LRANGE list 0 10                 # Get range

# Set operations
SADD tags "redis" "cache"
SMEMBERS tags
SISMEMBER tags "redis"

# Sorted Set (for leaderboards)
ZADD leaderboard 100 "user1"
ZRANGE leaderboard 0 10 WITHSCORES

# Cache patterns
SETEX key 3600 value            # Set with expiry
SETNX key value                 # Set if not exists (lock)
GETEX key EX 3600                # Get and refresh TTL
```

# Q&A Session

---

Time remaining for questions and discussion

Common Questions:

## 1. Redis vs Memcached?

- Redis: Feature-rich, persistent, data structures
- Memcached: Simpler, pure cache, slightly faster

## 2. How to size cache?

- Calculate: items × size × overhead
- Start with 20% of hot data
- Monitor and adjust

## 3. Cache in serverless (AWS Lambda)?

- Use external cache (ElastiCache)
- Or global variables (limited)
- Or DAX for DynamoDB

## 4. Handling cache cluster failures?

- Graceful degradation (fall back to DB)
- Cache replication
- Circuit breaker pattern

---

Session End - Thank You! 🎉

#caching #performance #redis #architecture #session-notes