

# Cache - Discussion Topics

## Architecture & Design

1. When to use cache-aside vs read-through vs write-through?
  - Trade-offs between consistency and performance
  - Use cases for each pattern
2. How do you handle cache invalidation in a distributed system?
  - Event-driven invalidation
  - TTL strategies
  - Cache coherence protocols
3. What are the risks of over-caching?
  - Memory pressure
  - Stale data issues
  - Complexity in debugging

## Real-World Scenarios

4. Design a caching strategy for a social media feed
  - User-specific vs global cache
  - Freshness requirements
  - Cache warming strategies
5. How would you prevent cache stampede during high traffic?
  - Request coalescing
  - Probabilistic early expiration
  - Lock-based approaches

## Performance & Optimization

6. How do you measure cache effectiveness?
  - Hit rate vs miss rate
  - Latency improvements
  - Cost-benefit analysis
7. What's your approach to cache sizing?
  - Working set estimation
  - Memory vs performance trade-offs
  - Monitoring and adjustment strategies

## Advanced Topics

8. Multi-level caching: When and how?
  - L1 (local) vs L2 (distributed) cache
  - Consistency challenges
  - Eviction coordination
9. How do you handle cache in a multi-region deployment?
  - Regional vs global cache
  - Replication strategies
  - Consistency models