Circuit Breaker & Resilience Patterns

Session Overview

ATTRIBUTE	DETAILS
Duration	60 minutes
Level	Intermediate to Advanced
Prerequisites	Microservices basics, distributed systems concepts

Agenda

| TIME	TOPIC
0-5 min	Introduction & Why Resilience Matters
5-20 min	Circuit Breaker Pattern Deep Dive
20-35 min	Bulkhead, Retry & Timeout Patterns
35-50 min	Fallback Strategies & Graceful Degradation
50-60 min	Practical Exercise & Discussion
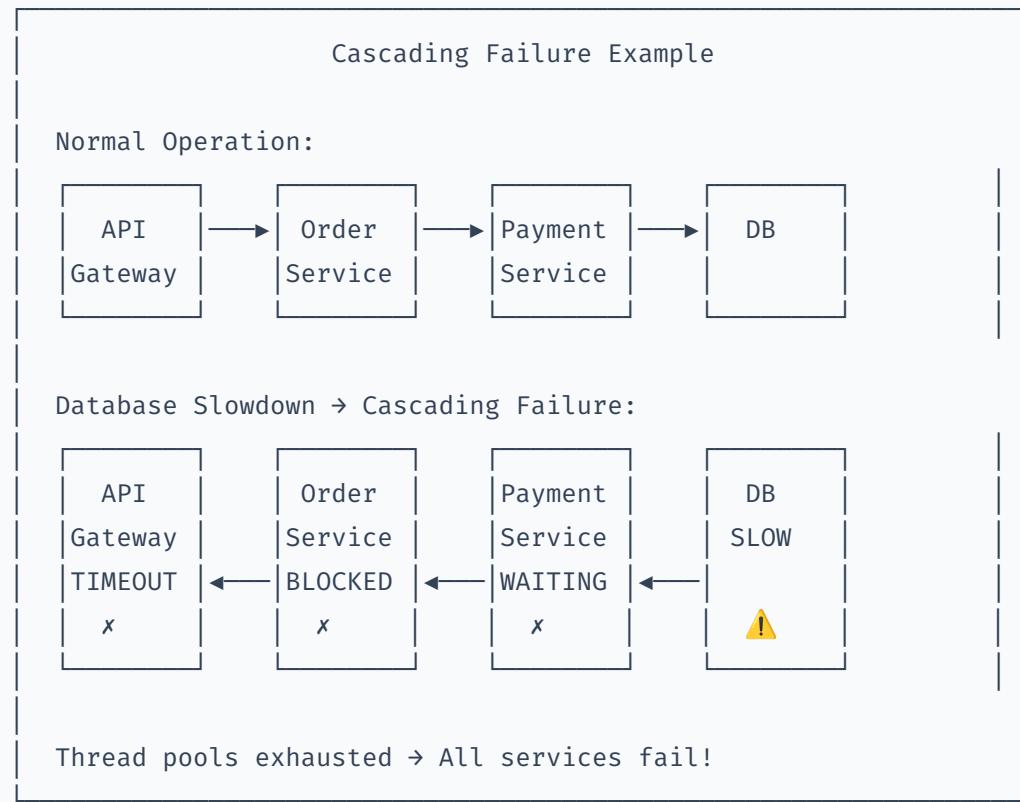
Learning Objectives

By the end of this session, you will be able to: - Understand the circuit breaker state machine and transitions - Implement bulkhead pattern for fault isolation - Design effective retry strategies with exponential backoff - Configure appropriate timeouts for different scenarios - Build fallback mechanisms for graceful degradation

1. Why Resilience Matters

The Cascading Failure Problem

```
┌─────────────────────────────────────────────────────────────┐
│                  Cascading Failure Example                   │
│                                                              │
│  Normal Operation:                                           │
│   ┌─────────┐    ┌─────────┐    ┌─────────┐    ┌─────────┐   │
│   │   API   │───▶│  Order  │───▶│ Payment │───▶│   DB    │   │
│   │ Gateway │    │ Service │    │ Service │    │         │   │
│   └─────────┘    └─────────┘    └─────────┘    └─────────┘   │
│                                                              │
│  Database Slowdown → Cascading Failure:                      │
│   ┌─────────┐    ┌─────────┐    ┌─────────┐    ┌─────────┐   │
│   │   API   │    │  Order  │    │ Payment │    │   DB    │   │
│   │ Gateway │    │ Service │    │ Service │    │  SLOW   │   │
│   │ TIMEOUT │◀───│ BLOCKED │◀───│ WAITING │◀───│         │   │
│   │    ✗    │    │    ✗    │    │    ✗    │    │   ⚠️    │   │
│   └─────────┘    └─────────┘    └─────────┘    └─────────┘   │
│                                                              │
│  Thread pools exhausted → All services fail!                 │
└─────────────────────────────────────────────────────────────┘
```
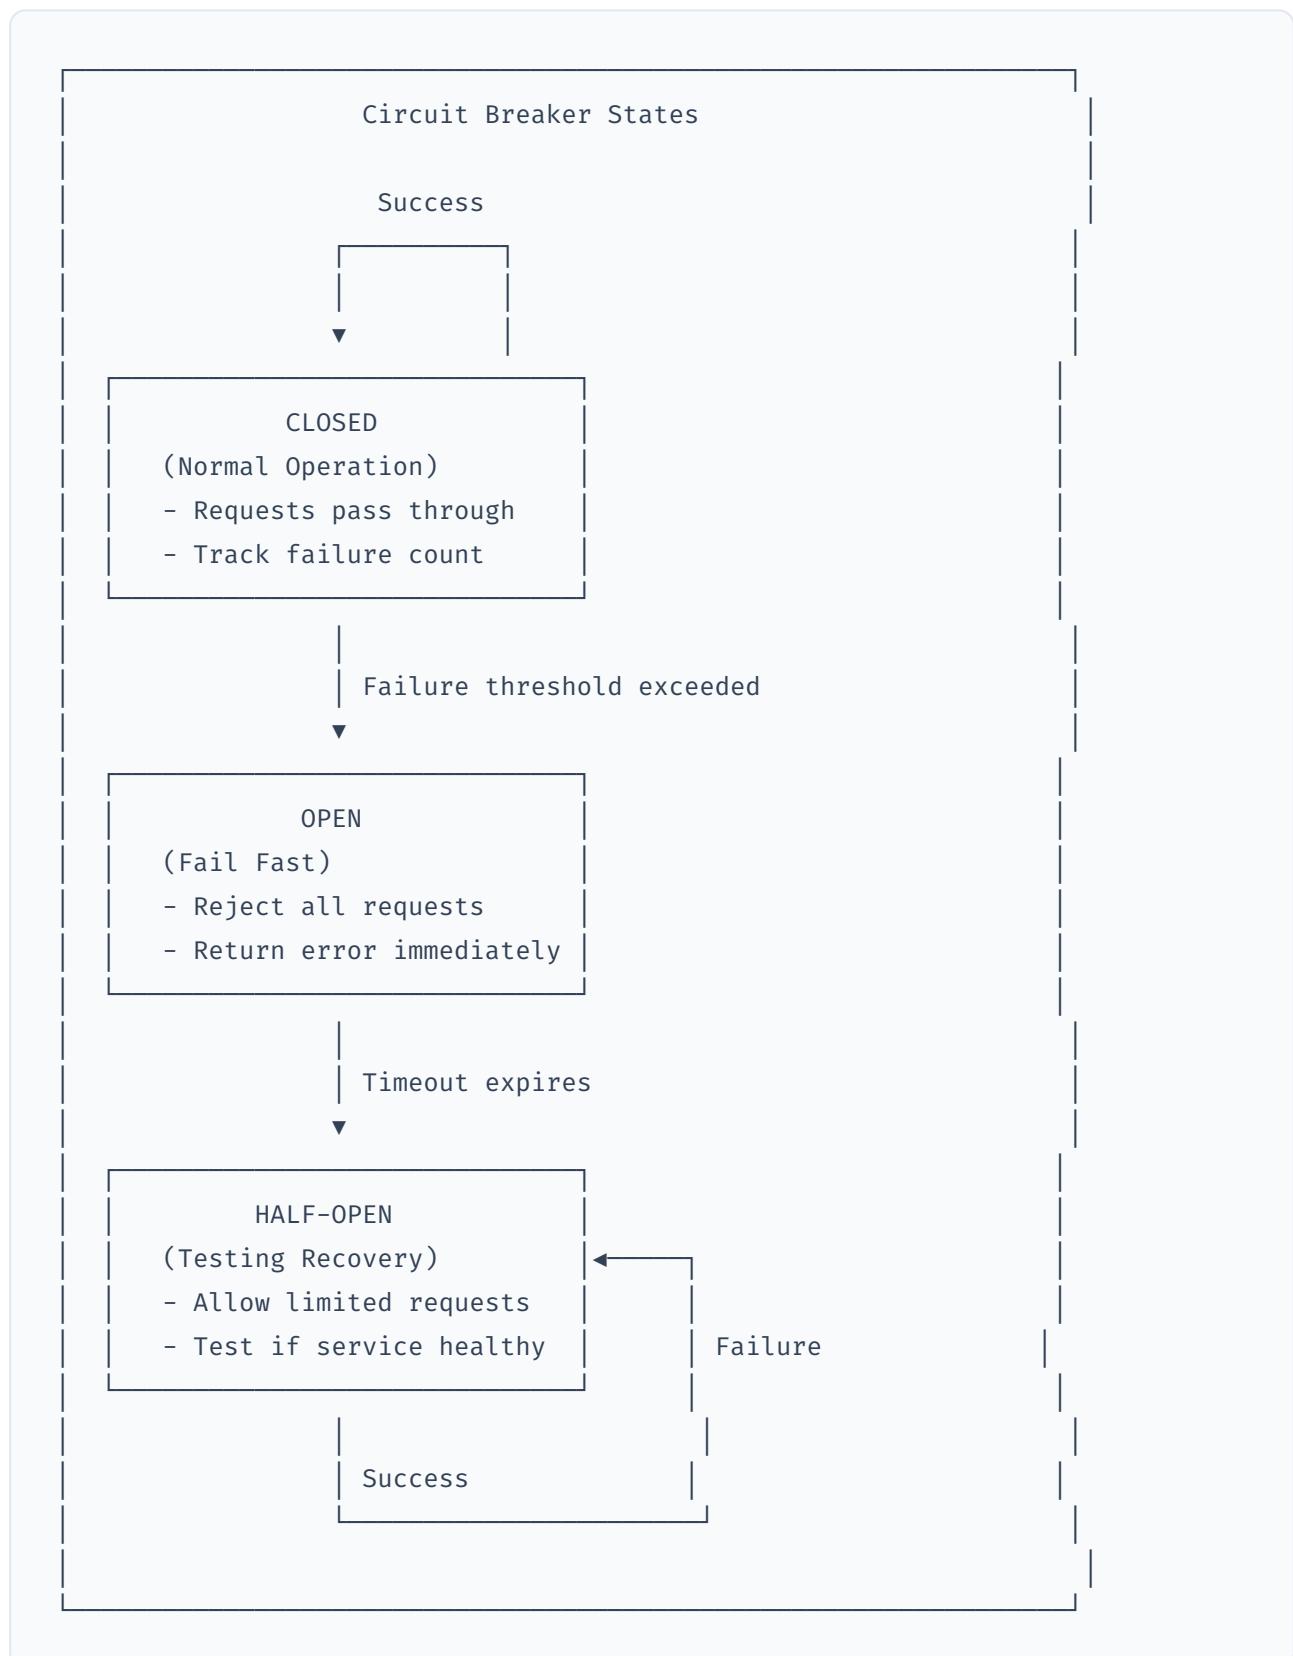
Resilience Patterns Overview

PATTERN	PURPOSE	WHEN TO USE
Circuit Breaker	Fail fast, prevent cascade	Downstream service failures
Bulkhead	Isolate failures	Resource contention
Retry	Handle transient failures	Network glitches, timeouts
Timeout	Bound wait time	Slow dependencies
Fallback	Graceful degradation	Any failure scenario

2. Circuit Breaker Pattern

State Machine

```
                      Circuit Breaker States

                          Success

                    ┌──────────┐
                    │          │
                    ▼          │
            ┌───────────────────┐
            │     CLOSED        │
            │ (Normal Operation)│
            │ - Requests pass through
            │ - Track failure count
            └───────────────────┘

                    │
                    │ Failure threshold exceeded
                    ▼
            ┌───────────────────┐
            │     OPEN          │
            │ (Fail Fast)       │
            │ - Reject all requests
            │ - Return error immediately
            └───────────────────┘

                    │
                    │ Timeout expires
                    ▼
            ┌───────────────────┐
            │   HALF-OPEN       │
            │ (Testing Recovery)│◄───┐
            │ - Allow limited requests │
            │ - Test if service healthy │  Failure
            └───────────────────┘    │
                    │                 │
                    │ Success         │
                    └─────────────────┘
```

Circuit Breaker Implementation

```java
public class CircuitBreaker {
    private final String name;
    private final int failureThreshold;
    private final long openTimeoutMs;
    private final int halfOpenMaxCalls;

    private CircuitState state = CircuitState.CLOSED;
    private int failureCount = 0;
    private int successCount = 0;
    private long lastFailureTime = 0;
    private int halfOpenCallCount = 0;

    public CircuitBreaker(String name, int failureThreshold,
                          long openTimeoutMs, int halfOpenMaxCalls) {
        this.name = name;
        this.failureThreshold = failureThreshold;
        this.openTimeoutMs = openTimeoutMs;
        this.halfOpenMaxCalls = halfOpenMaxCalls;
    }

    public <T> T execute(Supplier<T> action, Supplier<T> fallback) {
        if (!allowRequest()) {
            return fallback.get();
        }

        try {
            T result = action.get();
            recordSuccess();
            return result;
        } catch (Exception e) {
            recordFailure();
            return fallback.get();
        }
    }

    private synchronized boolean allowRequest() {
        switch (state) {
            case CLOSED:
                return true;

            case OPEN:
                if (System.currentTimeMillis() - lastFailureTime >
```

```java
      openTimeoutMs) {
                        transitionTo(CircuitState.HALF_OPEN);
                        return true;
                    }
                    return false;

                case HALF_OPEN:
                    if (halfOpenCallCount < halfOpenMaxCalls) {
                        halfOpenCallCount++;
                        return true;
                    }
                    return false;

                default:
                    return false;
            }
        }

        private synchronized void recordSuccess() {
            switch (state) {
                case CLOSED:
                    failureCount = 0;
                    break;

                case HALF_OPEN:
                    successCount++;
                    if (successCount ≥ halfOpenMaxCalls) {
                        transitionTo(CircuitState.CLOSED);
                    }
                    break;
            }
        }

        private synchronized void recordFailure() {
            lastFailureTime = System.currentTimeMillis();

            switch (state) {
                case CLOSED:
                    failureCount++;
                    if (failureCount ≥ failureThreshold) {
                        transitionTo(CircuitState.OPEN);
                    }
                    break;
```

```java
                case HALF_OPEN:
                    transitionTo(CircuitState.OPEN);
                    break;
            }
        }

        private void transitionTo(CircuitState newState) {
            CircuitState oldState = this.state;
            this.state = newState;
            this.failureCount = 0;
            this.successCount = 0;
            this.halfOpenCallCount = 0;

            log.info("Circuit breaker '{}' transitioned: {} → {}",
                    name, oldState, newState);
        }
    }

    enum CircuitState {
        CLOSED, OPEN, HALF_OPEN
    }
```

Using Resilience4j

```java
// Configuration
CircuitBreakerConfig config = CircuitBreakerConfig.custom()
    .failureRateThreshold(50)                    // 50% failure rate
    .waitDurationInOpenState(Duration.ofSeconds(30))
    .slidingWindowType(SlidingWindowType.COUNT_BASED)
    .slidingWindowSize(10)                       // Last 10 calls
    .minimumNumberOfCalls(5)                     // Min calls before
calculating
    .permittedNumberOfCallsInHalfOpenState(3)
    .automaticTransitionFromOpenToHalfOpenEnabled(true)
    .build();

CircuitBreaker circuitBreaker = CircuitBreaker.of("paymentService", config);

// Usage with decorator
Supplier<Payment> decoratedSupplier = CircuitBreaker
    .decorateSupplier(circuitBreaker, () → paymentService.process(order));

Try<Payment> result = Try.ofSupplier(decoratedSupplier)
    .recover(CallNotPermittedException.class, e → fallbackPayment());
```
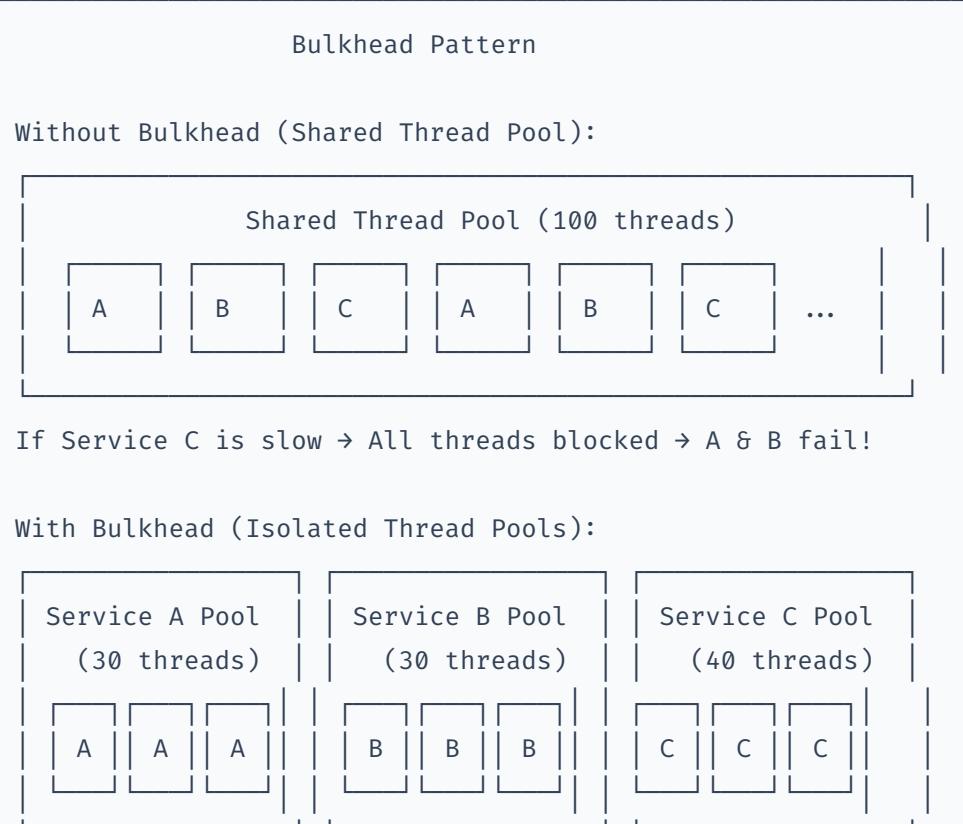
3. Bulkhead Pattern

Thread Pool Isolation

```
┌─────────────────────────────────────────────────────────────┐
│                      Bulkhead Pattern                        │
│                                                              │
│  Without Bulkhead (Shared Thread Pool):                      │
│  ┌────────────────────────────────────────────────────┐     │
│  │            Shared Thread Pool (100 threads)         │     │
│  │  ┌─────┐ ┌─────┐ ┌─────┐ ┌─────┐ ┌─────┐ ┌─────┐    │     │
│  │  │  A  │ │  B  │ │  C  │ │  A  │ │  B  │ │  C  │ ... │     │
│  │  └─────┘ └─────┘ └─────┘ └─────┘ └─────┘ └─────┘    │     │
│  └────────────────────────────────────────────────────┘     │
│  If Service C is slow → All threads blocked → A & B fail!    │
│                                                              │
│  With Bulkhead (Isolated Thread Pools):                      │
│  ┌──────────────┐ ┌──────────────┐ ┌──────────────┐          │
│  │ Service A Pool│ │ Service B Pool│ │ Service C Pool│         │
│  │  (30 threads) │ │  (30 threads) │ │  (40 threads) │         │
│  │ ┌──┐┌──┐┌──┐ │ │ ┌──┐┌──┐┌──┐ │ │ ┌──┐┌──┐┌──┐ │          │
│  │ │A ││A ││A │ │ │ │B ││B ││B │ │ │ │C ││C ││C │ │          │
│  │ └──┘└──┘└──┘ │ │ └──┘└──┘└──┘ │ │ └──┘└──┘└──┘ │          │
│  └──────────────┘ └──────────────┘ └──────────────┘          │
│  If Service C is slow → Only C's pool affected → A & B OK!   │
└─────────────────────────────────────────────────────────────┘
```

Bulkhead Implementation

```java
// Resilience4j Bulkhead Configuration
BulkheadConfig bulkheadConfig = BulkheadConfig.custom()
    .maxConcurrentCalls(25)          // Max concurrent calls
    .maxWaitDuration(Duration.ofMillis(500))  // Wait time for permit
    .build();

Bulkhead bulkhead = Bulkhead.of("paymentService", bulkheadConfig);

// Thread Pool Bulkhead (for async operations)
ThreadPoolBulkheadConfig threadPoolConfig = ThreadPoolBulkheadConfig.custom()
    .maxThreadPoolSize(10)
    .coreThreadPoolSize(5)
    .queueCapacity(100)
    .keepAliveDuration(Duration.ofMillis(100))
    .build();

ThreadPoolBulkhead threadPoolBulkhead =
    ThreadPoolBulkhead.of("paymentService", threadPoolConfig);

// Usage
Supplier<Payment> decoratedSupplier = Bulkhead
    .decorateSupplier(bulkhead, () → paymentService.process(order));
```

Semaphore-Based Bulkhead
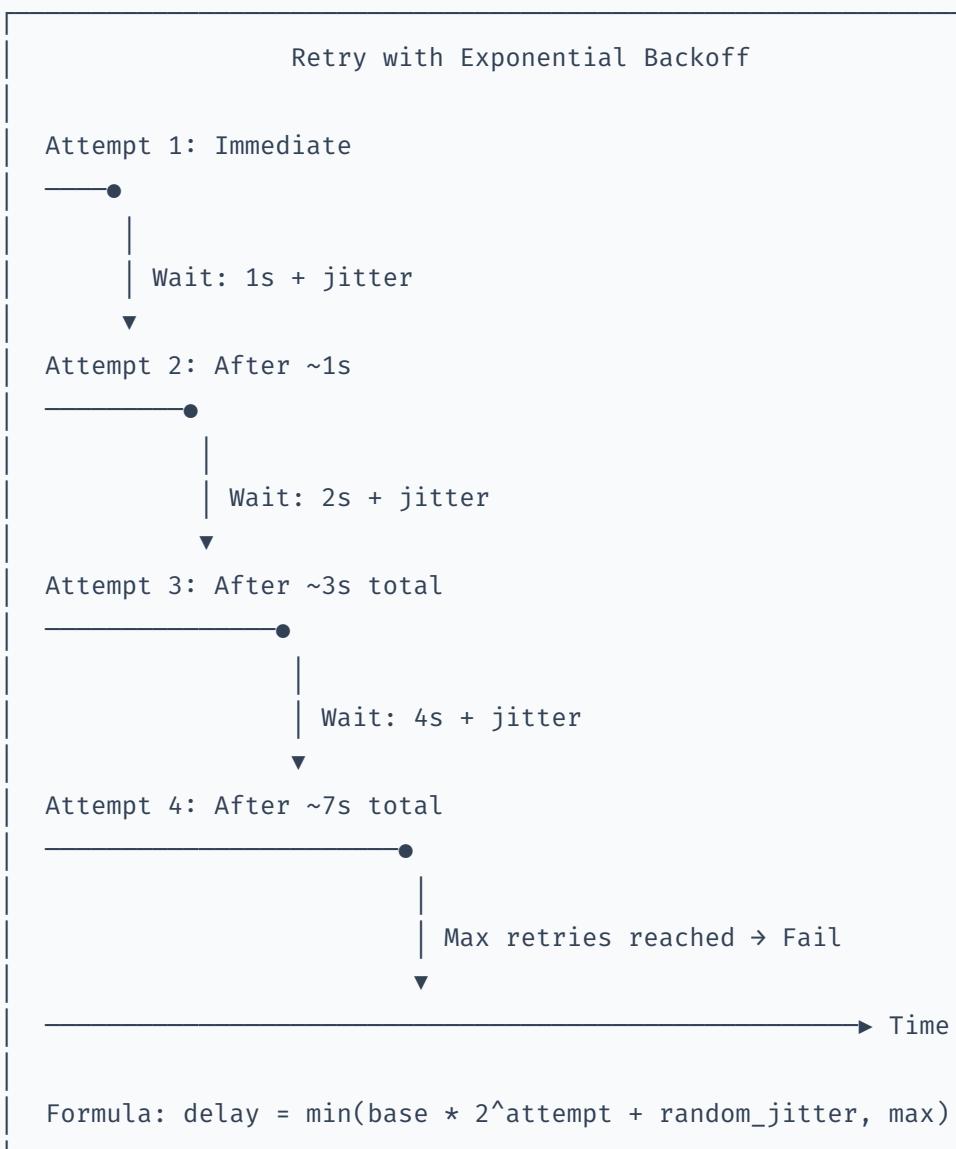
```java
public class SemaphoreBulkhead {
    private final Semaphore semaphore;
    private final String name;
    private final long maxWaitMs;

    public SemaphoreBulkhead(String name, int maxConcurrent, long maxWaitMs) {
        this.name = name;
        this.semaphore = new Semaphore(maxConcurrent);
        this.maxWaitMs = maxWaitMs;
    }

    public <T> T execute(Supplier<T> action) throws BulkheadFullException {
        boolean acquired = false;
        try {
            acquired = semaphore.tryAcquire(maxWaitMs,
TimeUnit.MILLISECONDS);
            if (!acquired) {
                throw new BulkheadFullException(
                    "Bulkhead '" + name + "' is full"
                );
            }
            return action.get();
        } catch (InterruptedException e) {
            Thread.currentThread().interrupt();
            throw new BulkheadFullException("Interrupted waiting for
bulkhead");
        } finally {
            if (acquired) {
                semaphore.release();
            }
        }
    }
}
```
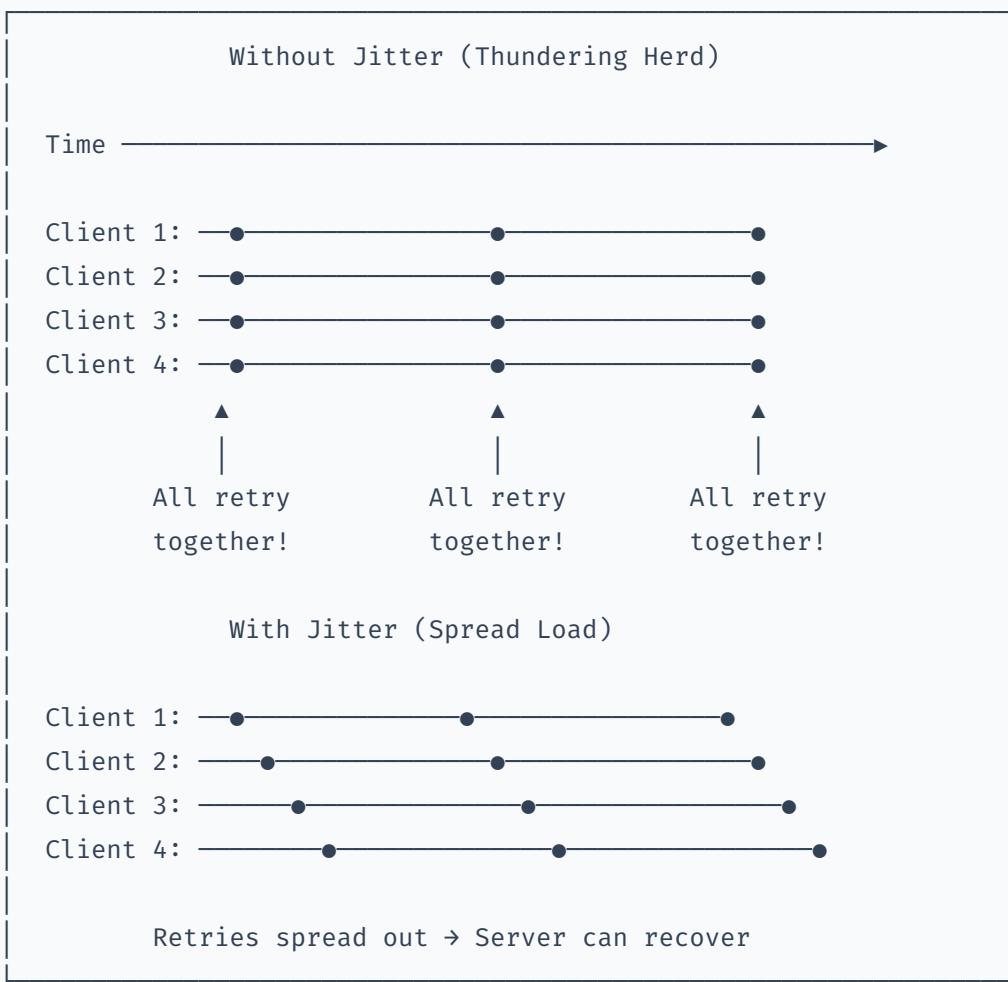
4. Retry Pattern

Exponential Backoff with Jitter

```
                    Retry with Exponential Backoff

  Attempt 1: Immediate
  ────────●
          │
          │  Wait: 1s + jitter
          ▼
  Attempt 2: After ~1s
  ──────────●
            │
            │  Wait: 2s + jitter
            ▼
  Attempt 3: After ~3s total
  ───────────────●
                 │
                 │  Wait: 4s + jitter
                 ▼
  Attempt 4: After ~7s total
  ──────────────────●
                    │
                    │  Max retries reached → Fail
                    ▼
  ──────────────────────────────────────────► Time

  Formula: delay = min(base * 2^attempt + random_jitter, max)
```

Why Jitter Matters

```
┌──────────────────────────────────────────────────────────┐
│              Without Jitter (Thundering Herd)            │
│                                                          │
│  Time ───────────────────────────────────────────►      │
│                                                          │
│  Client 1: ──●─────────────●─────────────●              │
│  Client 2: ──●─────────────●─────────────●              │
│  Client 3: ──●─────────────●─────────────●              │
│  Client 4: ──●─────────────●─────────────●              │
│             ▲             ▲             ▲               │
│             │             │             │               │
│           All retry     All retry     All retry         │
│           together!     together!     together!         │
│                                                          │
│              With Jitter (Spread Load)                  │
│                                                          │
│  Client 1: ──●─────────────●─────────────●              │
│  Client 2: ───●────────────●─────────────●              │
│  Client 3: ────●───────────●──────────────●             │
│  Client 4: ──────●─────────●───────────────●            │
│                                                          │
│        Retries spread out → Server can recover          │
└──────────────────────────────────────────────────────────┘
```

Retry Implementation

```java
public class RetryWithBackoff {
    private final int maxRetries;
    private final long baseDelayMs;
    private final long maxDelayMs;
    private final Set<Class<? extends Exception>> retryableExceptions;

    public RetryWithBackoff(int maxRetries, long baseDelayMs, long maxDelayMs) {
        this.maxRetries = maxRetries;
        this.baseDelayMs = baseDelayMs;
        this.maxDelayMs = maxDelayMs;
        this.retryableExceptions = Set.of(
            IOException.class,
            TimeoutException.class,
            ServiceUnavailableException.class
        );
    }

    public <T> T execute(Supplier<T> action) throws Exception {
        Exception lastException = null;

        for (int attempt = 0; attempt <= maxRetries; attempt++) {
            try {
                return action.get();
            } catch (Exception e) {
                lastException = e;

                if (!isRetryable(e) || attempt == maxRetries) {
                    throw e;
                }

                long delay = calculateDelay(attempt);
                log.warn("Attempt {} failed, retrying in {}ms: {}",
                        attempt + 1, delay, e.getMessage());
                Thread.sleep(delay);
            }
        }

        throw lastException;
    }

    private long calculateDelay(int attempt) {
```

```java
            // Exponential backoff: base * 2^attempt
            long exponentialDelay = baseDelayMs * (1L << attempt);

            // Add jitter: random value between 0 and delay
            long jitter = ThreadLocalRandom.current().nextLong(exponentialDelay /
2);

            // Cap at max delay
            return Math.min(exponentialDelay + jitter, maxDelayMs);
        }

        private boolean isRetryable(Exception e) {
            return retryableExceptions.stream()
                .anyMatch(clazz -> clazz.isInstance(e));
        }
    }

    // Using Resilience4j Retry
    RetryConfig retryConfig = RetryConfig.custom()
        .maxAttempts(3)
        .waitDuration(Duration.ofMillis(500))
        .intervalFunction(IntervalFunction.ofExponentialBackoff(
            Duration.ofMillis(500),  // Initial interval
            2.0,                      // Multiplier
            Duration.ofSeconds(10)   // Max interval
        ))
        .retryOnException(e -> e instanceof IOException)
        .ignoreExceptions(BusinessException.class)
        .build();

    Retry retry = Retry.of("paymentService", retryConfig);
```
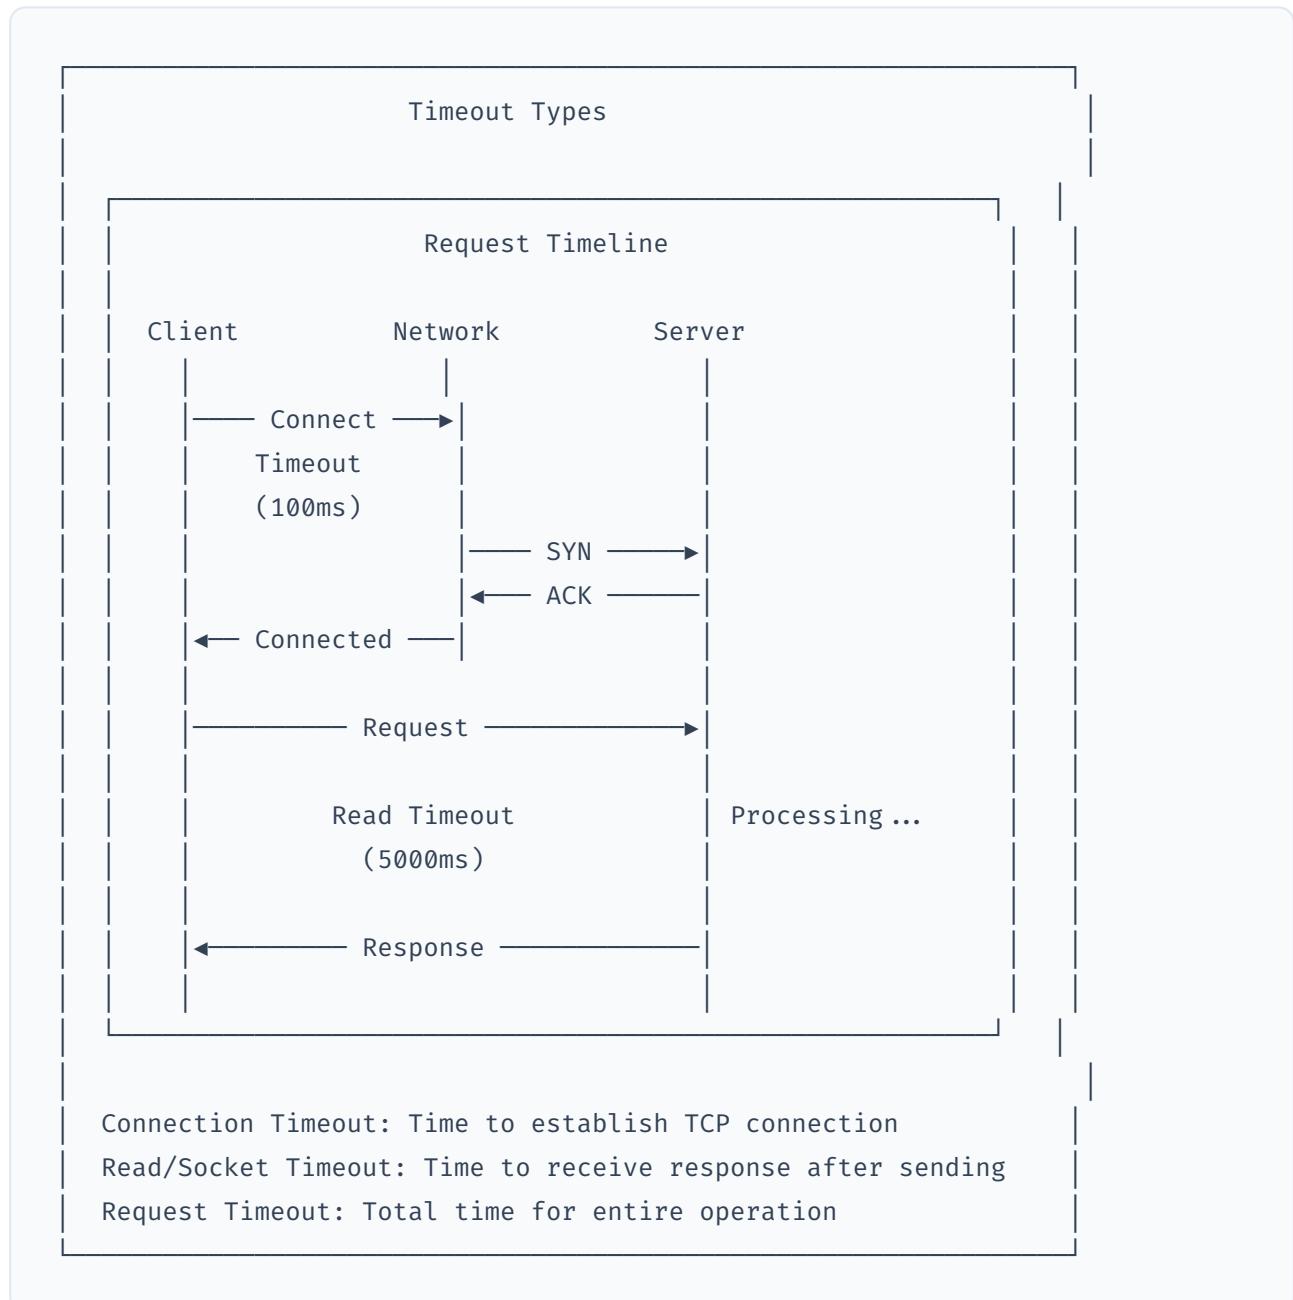
Retry Strategy Comparison

STRATEGY	FORMULA	USE CASE
Fixed Delay	`delay = constant`	Simple scenarios, low traffic
Linear Backoff	`delay = base * attempt`	Gradual increase
Exponential Backoff	`delay = base * 2^attempt`	Standard for distributed systems
Exponential + Jitter	`delay = base * 2^attempt + random`	High concurrency scenarios
Decorrelated Jitter	`delay = random(base, prev_delay * 3)`	AWS recommended approach

5. Timeout Pattern

Types of Timeouts

```
┌─────────────────────────────────────────────────────────┐
│                    Timeout Types                          │
│                                                           │
│   ┌─────────────────────────────────────────────────┐    │
│   │                 Request Timeline                  │    │
│   │                                                   │    │
│   │   Client          Network          Server         │    │
│   │      │                │                │          │    │
│   │      ├─── Connect ───▶│                │          │    │
│   │      │   Timeout      │                │          │    │
│   │      │   (100ms)      │                │          │    │
│   │      │                ├──── SYN ────▶  │          │    │
│   │      │                │◀─── ACK ─────  │          │    │
│   │      │◀── Connected ──│                │          │    │
│   │      │                │                │          │    │
│   │      ├───────── Request ──────────────▶│          │    │
│   │      │                │                │          │    │
│   │      │      Read Timeout      Processing...       │    │
│   │      │         (5000ms)               │          │    │
│   │      │                │                │          │    │
│   │      │◀──────── Response ─────────────│          │    │
│   │      │                │                │          │    │
│   └─────────────────────────────────────────────────┘    │
│                                                           │
│   Connection Timeout: Time to establish TCP connection    │
│   Read/Socket Timeout: Time to receive response after sending │
│   Request Timeout: Total time for entire operation        │
└─────────────────────────────────────────────────────────┘
```

Timeout Configuration

```java
// HTTP Client Timeouts
HttpClient httpClient = HttpClient.newBuilder()
    .connectTimeout(Duration.ofSeconds(2))
    .build();

HttpRequest request = HttpRequest.newBuilder()
    .uri(URI.create("https://api.example.com/payment"))
    .timeout(Duration.ofSeconds(5))  // Request timeout
    .build();

// RestTemplate Configuration
@Bean
public RestTemplate restTemplate() {
    HttpComponentsClientHttpRequestFactory factory =
        new HttpComponentsClientHttpRequestFactory();
    factory.setConnectTimeout(2000);      // Connection timeout
    factory.setReadTimeout(5000);         // Read timeout

    return new RestTemplate(factory);
}

// WebClient Configuration
@Bean
public WebClient webClient() {
    HttpClient httpClient = HttpClient.create()
        .option(ChannelOption.CONNECT_TIMEOUT_MILLIS, 2000)
        .responseTimeout(Duration.ofSeconds(5))
        .doOnConnected(conn -> conn
            .addHandlerLast(new ReadTimeoutHandler(5))
            .addHandlerLast(new WriteTimeoutHandler(5)));

    return WebClient.builder()
        .clientConnector(new ReactorClientHttpConnector(httpClient))
        .build();
}

// Resilience4j TimeLimiter
TimeLimiterConfig timeLimiterConfig = TimeLimiterConfig.custom()
    .timeoutDuration(Duration.ofSeconds(3))
    .cancelRunningFuture(true)
    .build();
```
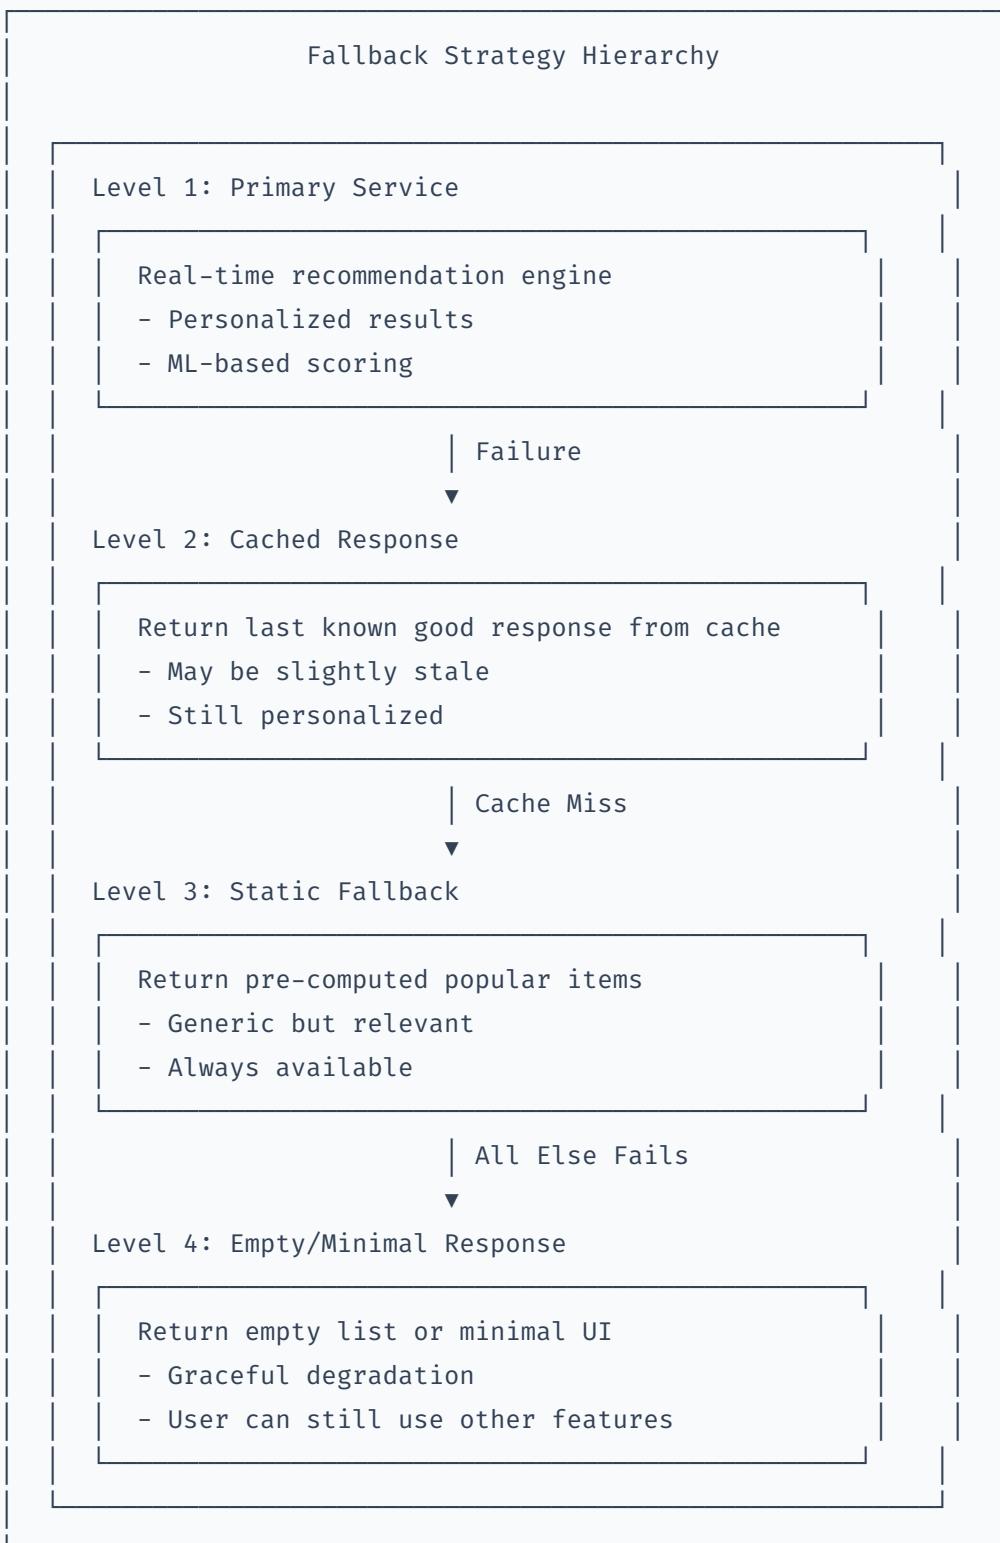
```
        TimeLimiter timeLimiter = TimeLimiter.of("paymentService",
timeLimiterConfig);
```

Timeout Guidelines

SERVICE TYPE	CONNECT TIMEOUT	READ TIMEOUT	NOTES
Internal APIs	100-500ms	1-5s	Low latency expected
External APIs	1-2s	5-30s	Higher variability
Database	1-2s	5-30s	Depends on query complexity
Cache (Redis)	100ms	500ms-1s	Should be fast
Message Queue	1s	5-10s	Depends on message size

6. Fallback Strategies

Graceful Degradation Patterns

```
                    Fallback Strategy Hierarchy

    ┌──────────────────────────────────────────────┐
    │  Level 1: Primary Service                      │
    │    ┌──────────────────────────────────────┐   │
    │    │  Real-time recommendation engine       │   │
    │    │  - Personalized results                │   │
    │    │  - ML-based scoring                    │   │
    │    └──────────────────────────────────────┘   │
    │                       │ Failure                │
    │                       ▼                        │
    │  Level 2: Cached Response                      │
    │    ┌──────────────────────────────────────┐   │
    │    │  Return last known good response from cache │
    │    │  - May be slightly stale               │   │
    │    │  - Still personalized                  │   │
    │    └──────────────────────────────────────┘   │
    │                       │ Cache Miss             │
    │                       ▼                        │
    │  Level 3: Static Fallback                      │
    │    ┌──────────────────────────────────────┐   │
    │    │  Return pre-computed popular items     │   │
    │    │  - Generic but relevant                │   │
    │    │  - Always available                    │   │
    │    └──────────────────────────────────────┘   │
    │                       │ All Else Fails         │
    │                       ▼                        │
    │  Level 4: Empty/Minimal Response               │
    │    ┌──────────────────────────────────────┐   │
    │    │  Return empty list or minimal UI       │   │
    │    │  - Graceful degradation                │   │
    │    │  - User can still use other features   │   │
    │    └──────────────────────────────────────┘   │
    └──────────────────────────────────────────────┘
```

Fallback Implementation

```java
@Service
public class ProductRecommendationService {

    private final RecommendationEngine recommendationEngine;
    private final CacheService cacheService;
    private final CircuitBreaker circuitBreaker;

    public List<Product> getRecommendations(String userId) {
        return circuitBreaker.execute(
            () -> getPrimaryRecommendations(userId),
            () -> getFallbackRecommendations(userId)
        );
    }

    private List<Product> getPrimaryRecommendations(String userId) {
        // Primary: Real-time ML-based recommendations
        List<Product> recommendations =
recommendationEngine.recommend(userId);

        // Cache for fallback
        cacheService.put("recommendations:" + userId, recommendations,
                        Duration.ofHours(1));

        return recommendations;
    }

    private List<Product> getFallbackRecommendations(String userId) {
        // Level 1: Try cached personalized recommendations
        Optional<List<Product>> cached = cacheService.get(
            "recommendations:" + userId,
            new TypeReference<List<Product>>() {}
        );

        if (cached.isPresent()) {
            log.info("Returning cached recommendations for user {}", userId);
            return cached.get();
        }

        // Level 2: Try category-based popular items
        String userCategory = getUserCategory(userId);
        Optional<List<Product>> categoryPopular = cacheService.get(
            "popular:" + userCategory,
```

```java
                new TypeReference<List<Product>>() {}
            );

            if (categoryPopular.isPresent()) {
                log.info("Returning category popular items for user {}", userId);
                return categoryPopular.get();
            }

            // Level 3: Return global popular items
            log.warn("Returning global fallback for user {}", userId);
            return getGlobalPopularItems();
        }

    @Cacheable("global-popular")
    public List<Product> getGlobalPopularItems() {
        // Pre-computed list of globally popular items
        return productRepository.findTopByOrderBySalesDesc(20);
    }
}
```

Combining All Patterns

```java
@Service
public class ResilientPaymentService {

    private final CircuitBreaker circuitBreaker;
    private final Bulkhead bulkhead;
    private final Retry retry;
    private final TimeLimiter timeLimiter;

    @PostConstruct
    public void init() {
        // Circuit Breaker: 50% failure rate opens circuit
        circuitBreaker = CircuitBreaker.of("payment",
CircuitBreakerConfig.custom()
                .failureRateThreshold(50)
                .waitDurationInOpenState(Duration.ofSeconds(30))
                .slidingWindowSize(10)
                .build());

        // Bulkhead: Max 20 concurrent calls
        bulkhead = Bulkhead.of("payment", BulkheadConfig.custom()
            .maxConcurrentCalls(20)
            .maxWaitDuration(Duration.ofMillis(500))
            .build());

        // Retry: 3 attempts with exponential backoff
        retry = Retry.of("payment", RetryConfig.custom()
            .maxAttempts(3)
            .intervalFunction(IntervalFunction.ofExponentialBackoff(
                Duration.ofMillis(100), 2.0))
            .retryExceptions(IOException.class, TimeoutException.class)
            .build());

        // Timeout: 5 second limit
        timeLimiter = TimeLimiter.of("payment", TimeLimiterConfig.custom()
            .timeoutDuration(Duration.ofSeconds(5))
            .build());
    }

    public PaymentResult processPayment(PaymentRequest request) {
        // Compose decorators: Bulkhead → CircuitBreaker → Retry →
TimeLimiter
        Supplier<PaymentResult> decoratedSupplier = Decorators
```

```
                .ofSupplier(() → paymentGateway.process(request))
                .withBulkhead(bulkhead)
                .withCircuitBreaker(circuitBreaker)
                .withRetry(retry)
                .withTimeLimiter(timeLimiter, executorService)
                .withFallback(Arrays.asList(
                    TimeoutException.class,
                    BulkheadFullException.class,
                    CallNotPermittedException.class
                ), e → handleFallback(request, e))
                .decorate();

        return Try.ofSupplier(decoratedSupplier)
            .getOrElseThrow(e → new PaymentException("Payment failed", e));
    }

    private PaymentResult handleFallback(PaymentRequest request, Throwable e)
{
        log.error("Payment failed, executing fallback: {}", e.getMessage());

        // Queue for async processing
        paymentQueue.enqueue(request);

        return PaymentResult.builder()
            .status(PaymentStatus.PENDING)
            .message("Payment queued for processing")
            .build();
    }
}
```

Key Takeaways

1. **Circuit Breaker** prevents cascading failures by failing fast when a service is unhealthy

2. **Bulkhead** isolates failures by limiting concurrent access to resources

3. **Retry with backoff** handles transient failures while avoiding thundering herd

4. **Timeouts** bound wait times and prevent resource exhaustion

5. **Fallback strategies** provide graceful degradation when primary services fail

6. **Combine patterns** for comprehensive resilience (order matters: Bulkhead → Circuit Breaker → Retry → Timeout)

Practical Exercise

Scenario: Build a Resilient Product Service

Design and implement a resilient product service that:

1. Fetches product details from an external inventory API

2. Gets pricing from a pricing service

3. Retrieves reviews from a review service

Requirements: - Circuit breaker with 50% failure threshold - Bulkhead limiting to 10 concurrent calls per downstream service - Retry with exponential backoff (max 3 attempts) - 3-second timeout for each service call - Fallback to cached data when services fail

Starter Code:

```java
@Service
public class ProductService {

    // TODO: Add resilience patterns

    public ProductDetails getProduct(String productId) {
        // Fetch from multiple services
        Inventory inventory = inventoryService.getInventory(productId);
        Price price = pricingService.getPrice(productId);
        List<Review> reviews = reviewService.getReviews(productId);

        return ProductDetails.builder()
            .productId(productId)
            .inventory(inventory)
            .price(price)
            .reviews(reviews)
            .build();
    }
}
```

Discussion Questions: 1. How would you handle partial failures (e.g., reviews fail but inventory succeeds)? 2. What metrics would you monitor to tune the circuit breaker thresholds? 3. How would you test the resilience patterns in a staging environment?