

Test creaeting pdf from markdown

# Consistent Hashing - 1 Hour Session

---

**Duration:** 60 minutes **Level:** Intermediate

---



## Session Agenda

---

- Introduction to Consistent Hashing (10 min)
  - Hash Ring Fundamentals (15 min)
  - Virtual Nodes (10 min)
  - Rebalancing & Data Migration (10 min)
  - Implementation & Real-World Applications (15 min)
- 

## 🎯 Learning Objectives

---

By the end of this session, you will understand:

- Why traditional hashing fails in distributed systems
- How consistent hashing solves the redistribution problem
- The hash ring concept and key placement
- Virtual nodes and their benefits
- How to handle node additions and removals
- Real-world applications in distributed databases and caches

---

# 1. Introduction to Consistent Hashing (10 min)

## The Problem with Traditional Hashing

### Simple Modulo Hashing:

```
server_index = hash(key) % number_of_servers
```

Example with 4 servers:

```
hash("user:123") = 12345  
server = 12345 % 4 = 1 → Server 1
```

### What happens when we add/remove a server?

```
graph TD; subgraph Before ["Before: 4 servers"]; K1[hash('user:123') % 4 = 1] --> S1[Server 1]; K2[hash('user:456') % 4 = 0] --> S0[Server 0]; K3[hash('user:789') % 4 = 3] --> S3[Server 3]; end; subgraph After ["After: 5 servers"]; K1a[hash('user:123') % 5 = 0] --> S0a[Server 0]; K2a[hash('user:456') % 5 = 0] --> S0b[Server 0]; K3a[hash('user:789') % 5 = 4] --> S4[Server 4]; end; Before --> After;
```

**Result: ~80% of keys need to be remapped! ⚡**

**⚠ Warning** With  $N$  servers, adding or removing one server causes  $(N-1)/N$  keys to be remapped. For 100 servers, that's 99% of all data!

## Why This Matters

**Impact on Distributed Systems:** - 🔥 **Cache Invalidations:** Massive cache misses during scaling - 📊 **Database Sharding:** Expensive data migration - ⏳ **Downtime:** System unavailable during rebalancing - 💰 **Cost:** Network bandwidth and compute for data movement

## The Solution: Consistent Hashing

**Key Insight:** Instead of mapping keys to servers directly, map both keys AND servers to a fixed hash space (ring).



**Benefits:** - Only K/N keys remapped when adding/removing nodes (K = total keys, N = nodes) - Minimal data movement during scaling - Predictable load distribution

## 2. Hash Ring Fundamentals (15 min)

### The Hash Ring Concept

Creating the Ring:

```

flowchart TB
    subgraph Ring ["Hash Ring (0 to 2^32 - 1)"]
        direction TB
        P0["0°"] — P90["90°"]
        P90 — P180["180°"]
        P180 — P270["270°"]
        P270 — P0
    end

    Formula["Position = hash(identifier) % ring_size"]

```

## Placing Servers on the Ring

```

import hashlib

def hash_to_ring(key, ring_size=2**32):
    """Hash a key to a position on the ring"""
    hash_value = int(hashlib.md5(key.encode()).hexdigest(), 16)
    return hash_value % ring_size

# Place servers on the ring
servers = ["server-A", "server-B", "server-C", "server-D"]
server_positions = {s: hash_to_ring(s) for s in servers}

# Example positions (simplified to 0-360 for visualization)
# server-A: 45°
# server-B: 120°
# server-C: 200°
# server-D: 300°

```

### Visual Representation:

```

flowchart TB
    subgraph HashRing["Hash Ring"]
        direction TB
        P0["0°"]
        A["server-A (45°)"]
        B["server-B (120°)"]
        C["server-C (200°)"]
        D["server-D (300°)"]

        P0 --> A --> B --> C --> D --> P0
    end

```

## Key Placement: Finding the Right Server

**Rule:** A key is assigned to the first server encountered when moving clockwise from the key's position.

```

def find_server(key, server_positions):
    """Find the server responsible for a key"""
    key_position = hash_to_ring(key)

    # Sort servers by position
    sorted_servers = sorted(server_positions.items(), key=lambda x: x[1])

    # Find first server clockwise from key position
    for server, position in sorted_servers:
        if position >= key_position:
            return server

    # Wrap around to first server
    return sorted_servers[0][0]

```

### Example:

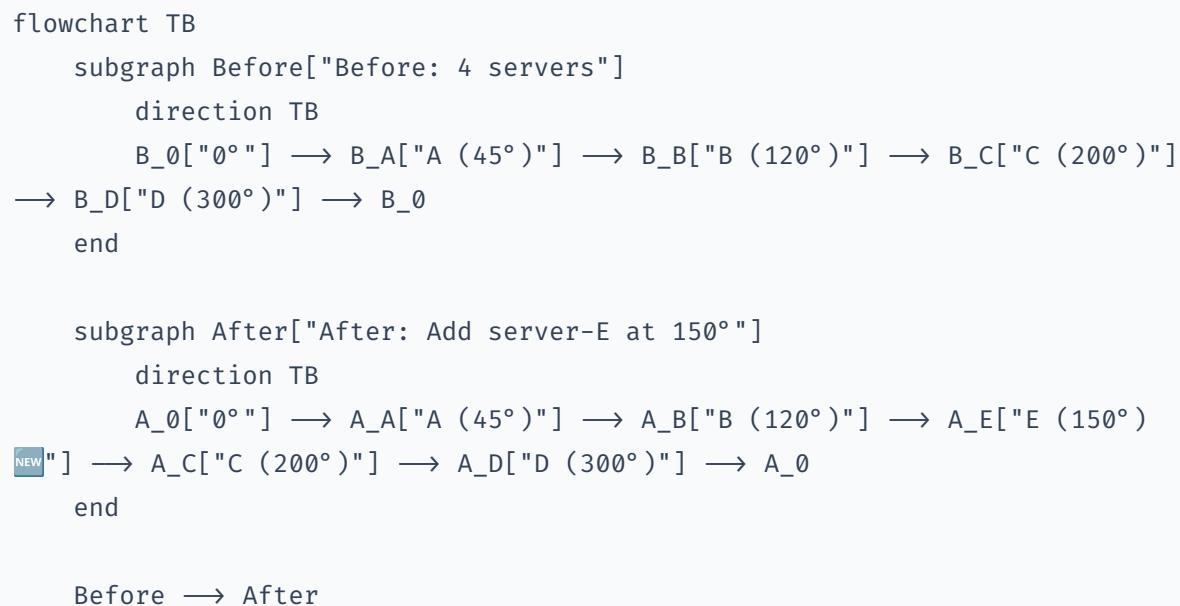
```
Key "user:123" hashes to position 80°
```

Ring positions:

- server-A: 45°
- server-B: 120° ← First server clockwise from 80°
- server-C: 200°
- server-D: 300°

Result: "user:123" → server-B

## Adding a New Server



**Only keys between 120° and 150° move from C to E!**

**Impact Analysis:**

```

# Keys affected when adding server-E at 150°
# Only keys in range (120°, 150°] move from server-C to server-E

# Before: Keys 121°-200° → server-C
# After:   Keys 121°-150° → server-E
#           Keys 151°-200° → server-C

# Percentage moved: (150-120)/(360) ≈ 8.3% of total keys
# Much better than 80% with modulo hashing!

```

## Removing a Server

Remove server-B (120°):

Before:

- Keys 46°-120° → server-B

After:

- Keys 46°-120° → server-E (next clockwise server)

Only server-B's keys need to move!

## 3. Virtual Nodes (10 min)

### The Load Imbalance Problem

With few physical nodes, distribution can be uneven:

```
4 servers, random positions:  
- server-A: 10° (covers 10° of ring = 2.8%)  
- server-B: 50° (covers 40° of ring = 11.1%)  
- server-C: 300° (covers 250° of ring = 69.4%)  
- server-D: 350° (covers 50° of ring = 13.9%)
```

server-C handles 69% of all keys! 💥

## Solution: Virtual Nodes (VNodes)

Idea: Each physical server gets multiple positions on the ring.

```
def create_virtual_nodes(server, num_vnodes=150):  
    """Create multiple positions for a single server"""  
    vnodes = []  
    for i in range(num_vnodes):  
        vnode_key = f"{server}#vnode{i}"  
        position = hash_to_ring(vnode_key)  
        vnodes[position] = server  
    return vnodes  
  
# Example: server-A with 3 virtual nodes  
# server-A#vnode0 → position 45°  
# server-A#vnode1 → position 180°  
# server-A#vnode2 → position 290°
```

## Visual with Virtual Nodes:

```

flowchart TB
    subgraph VNodeRing["Ring with Virtual Nodes (Physical: A, B)"]
        direction TB
        P0["P0 [0°]"]
        A0["A#0 (45°)"]
        B0["B#0 (90°)"]
        A1["A#1 (150°)"]
        B1["B#1 (200°)"]
        B2["B#2 (250°)"]
        A2["A#2 (320°)"]

        P0 --> A0 --> B0 --> A1 --> B1 --> B2 --> A2 --> P0
    end

    style A0 fill:#4CAF50
    style A1 fill:#4CAF50
    style A2 fill:#4CAF50
    style B0 fill:#2196F3
    style B1 fill:#2196F3
    style B2 fill:#2196F3

```

**Now load is distributed more evenly!**

## Benefits of Virtual Nodes

ASPECT	WITHOUT VNODLES	WITH VNODLES
Load Distribution	Uneven	Even
Adding Node	Affects 1 range	Affects many small ranges
Heterogeneous Hardware	Not supported	Supported (more vnodes for powerful servers)
Failure Impact	Large chunk unavailable	Small chunks distributed

## Handling Heterogeneous Hardware

```
# More powerful servers get more virtual nodes
server_vnodes = {
    "server-A": 100,      # Standard server
    "server-B": 100,      # Standard server
    "server-C": 200,      # 2x capacity server
    "server-D": 50,       # Half capacity server
}

# server-C handles ~2x the load of A or B
# server-D handles ~0.5x the load of A or B
```

## Choosing the Number of Virtual Nodes

Factors to consider:

- More vnodes = better distribution, but more memory for ring
- Typical range: 100-200 vnodes per physical node
- Amazon DynamoDB uses 256 vnodes per node

Memory calculation:

- 100 physical nodes × 150 vnodes = 15,000 ring entries
- Each entry: ~100 bytes (position + server reference)
- Total: ~1.5 MB (negligible)

## 4. Rebalancing & Data Migration (10 min)

### When Rebalancing Occurs

1. **Node Addition:** New server joins the cluster
2. **Node Removal:** Server leaves (planned or failure)
3. **Capacity Change:** Adjusting virtual node count

## Rebalancing Strategy

### Adding a Node:

```
def add_node(ring, new_server, num_vnodes=150):
    """Add a new server to the ring"""
    keys_to_migrate = {}

    for i in range(num_vnodes):
        vnode_key = f"{new_server}#vnode{i}"
        new_position = hash_to_ring(vnode_key)

        # Find the server that currently owns this position
        current_owner = find_server_at_position(ring, new_position)

        # Keys between previous position and new position move to new server
        prev_position = find_previous_position(ring, new_position)
        keys_to_migrate[new_position] = {
            'from': current_owner,
            'range': (prev_position, new_position)
        }

        # Add new vnode to ring
        ring[new_position] = new_server

    return keys_to_migrate
```

### Migration Process:

Step 1: Add new node to ring (metadata only)  
Step 2: New node starts accepting writes for its range  
Step 3: Background migration of existing keys  
Step 4: Old node stops serving migrated keys

#### Timeline:

t=0 [Add to ring]  
t=1 [Accept new writes]  
t=2-10 [Background migration]  
t=11 [Complete]

## Handling Node Failures

### Graceful Removal:

```
def remove_node_graceful(ring, server):
    """Gracefully remove a server"""
    # 1. Stop accepting new writes
    server.stop_writes()

    # 2. Migrate data to successor nodes
    for vnode_position in get_vnodes(ring, server):
        successor = find_next_server(ring, vnode_position)
        migrate_data(server, successor, vnode_position)

    # 3. Remove from ring
    remove_vnodes(ring, server)
```

### Sudden Failure:

When a node fails unexpectedly:

1. Detect failure (heartbeat timeout)
2. Remove from ring immediately
3. Successor nodes now responsible for failed node's range
4. If replication exists, replicas serve requests
5. Re-replicate data to maintain replication factor

## Data Replication with Consistent Hashing

```
flowchart TB
    subgraph Replication ["Replication Factor: 3"]
        direction TB
        Key["Key 'user:123' at 80°"]

        subgraph Ring ["Hash Ring"]
            A["A (45°)"]
            B["B (120°) <br/> ● Primary"]
            E["E (150°) <br/> ● Replica 1"]
            C["C (200°) <br/> ○ Replica 2"]
            D["D (300°)"]
        end

        A --> B --> E --> C --> D --> A
    end

    Key -.→|"clockwise"| B
end

style B fill:#2196F3
style E fill:#4CAF50
style C fill:#FFC107
```

**Strategy:** Replicate to N-1 successor nodes

## **5. Implementation & Real-World Applications (15 min)**

---

### **Complete Implementation**

```

import hashlib
from bisect import bisect_right

class ConsistentHashRing:
    def __init__(self, num_vnodes=150):
        self.num_vnodes = num_vnodes
        self.ring = {} # position → server
        self.sorted_positions = []
        self.servers = set()

    def _hash(self, key):
        """Generate hash for a key"""
        return int(hashlib.md5(key.encode()).hexdigest(), 16)

    def add_server(self, server):
        """Add a server with virtual nodes"""
        self.servers.add(server)
        for i in range(self.num_vnodes):
            vnode_key = f"{server}#vnode{i}"
            position = self._hash(vnode_key)
            self.ring[position] = server

        self.sorted_positions = sorted(self.ring.keys())

    def remove_server(self, server):
        """Remove a server and its virtual nodes"""
        self.servers.discard(server)
        positions_to_remove = [
            pos for pos, srv in self.ring.items() if srv == server
        ]
        for pos in positions_to_remove:
            del self.ring[pos]

        self.sorted_positions = sorted(self.ring.keys())

    def get_server(self, key):
        """Find the server responsible for a key"""
        if not self.ring:
            return None

        key_hash = self._hash(key)

```

```

        # Find first position ≥ key_hash
        idx = bisect_right(self.sorted_positions, key_hash)

        # Wrap around if necessary
        if idx == len(self.sorted_positions):
            idx = 0

        position = self.sorted_positions[idx]
        return self.ring[position]

    def get_servers_for_replication(self, key, replicas=3):
        """Get N servers for replication"""
        if not self.ring or replicas > len(self.servers):
            return list(self.servers)

        key_hash = self._hash(key)
        idx = bisect_right(self.sorted_positions, key_hash)

        servers = []
        seen = set()

        while len(servers) < replicas:
            if idx ≥ len(self.sorted_positions):
                idx = 0

            server = self.ring[self.sorted_positions[idx]]
            if server not in seen:
                servers.append(server)
                seen.add(server)

            idx += 1

        return servers

# Usage
ring = ConsistentHashRing(num_vnodes=150)
ring.add_server("server-A")
ring.add_server("server-B")
ring.add_server("server-C")

# Find server for a key
server = ring.get_server("user:12345")
print(f"user:12345 → {server}")

```

```
# Get replicas
replicas = ring.get_servers_for_replication("user:12345", replicas=3)
print(f"Replicas: {replicas}")
```

## Real-World Applications

### 1. Amazon DynamoDB

DynamoDB Partitioning:

- Uses consistent hashing for partition key distribution
- 256 virtual nodes per physical partition
- Automatic rebalancing during scaling

Partition Key → Hash → Partition

"user:123" → 0x7F ... → Partition-B

### 2. Apache Cassandra

Cassandra Token Ring:

- Each node owns a range of tokens
- Virtual nodes (vnodes) enabled by default
- Configurable replication factor

```
CREATE KEYSPACE my_keyspace
WITH replication = {
    'class': 'NetworkTopologyStrategy',
    'datacenter1': 3
};
```

### 3. Redis Cluster

Redis Cluster:

- 16384 hash slots (not continuous ring)
- Each node owns a subset of slots
- Manual or automatic slot migration

```
CLUSTER ADDSLOTS 0 1 2 3 ... 5460 # Node A  
CLUSTER ADDSLOTS 5461 ... 10922 # Node B  
CLUSTER ADDSLOTS 10923 ... 16383 # Node C
```

### 4. Content Delivery Networks (CDNs)

CDN Edge Server Selection:

- Hash(content\_url) → Edge server
- Ensures same content always served from same edge
- Maximizes cache hit rate

Request: GET /images/logo.png

Hash: 0x3F ...

Edge: edge-server-tokyo-2

## Load Balancing with Consistent Hashing

```
class ConsistentHashLoadBalancer:
    def __init__(self):
        self.ring = ConsistentHashRing(num_vnodes=100)
        self.server_health = {}

    def add_backend(self, server, weight=1):
        """Add backend server with optional weight"""
        # More vnodes for higher weight
        vnodes = 100 * weight
        self.ring.num_vnodes = vnodes
        self.ring.add_server(server)
        self.server_health[server] = True

    def get_backend(self, request_key):
        """Get backend for a request"""
        # Use session ID or client IP as key for sticky sessions
        server = self.ring.get_server(request_key)

        # Health check
        if not self.server_health.get(server, False):
            # Fallback to next healthy server
            return self._get_next_healthy(request_key)

    return server

    def _get_next_healthy(self, key):
        """Find next healthy server on the ring"""
        servers = self.ring.get_servers_for_replication(key,
replicas=len(self.ring.servers))
        for server in servers:
            if self.server_health.get(server, False):
                return server
        return None
```



# Common Pitfalls & Solutions

## 1. Hot Spots

**Problem:** Some keys are accessed much more frequently.

```
Popular user "celebrity:123" → server-B  
Server-B gets 100x more traffic than others
```

**Solutions:** - Add read replicas for hot keys - Cache hot keys at application layer - Split hot keys across multiple entries

## 2. Uneven Virtual Node Distribution

**Problem:** Random hashing can still create clusters.

**Solution:** Use deterministic placement or jump consistent hashing.

```
# Jump Consistent Hash (Google)  
def jump_consistent_hash(key, num_buckets):  
    b, j = -1, 0  
    while j < num_buckets:  
        b = j  
        key = ((key * 2862933555777941757) + 1) & 0xFFFFFFFFFFFFFF  
        j = int((b + 1) * (float(1 << 31) / float((key >> 33) + 1)))  
    return b
```

## 3. Ring Metadata Synchronization

**Problem:** All nodes need consistent view of the ring.

**Solutions:** - Gossip protocol (Cassandra) - Centralized coordinator (ZooKeeper) - Consensus algorithm (Raft)

## 4. Data Migration During Rebalancing

**Problem:** Large data movement affects performance.

**Solutions:** - Background migration with rate limiting - Dual-write during transition - Lazy migration (migrate on access)

## Comparison with Alternatives

APPROACH	REDISTRIBUTION	COMPLEXITY	USE CASE
Modulo Hash	~100% on change	Low	Static clusters
Consistent Hash	~K/N on change	Medium	Dynamic clusters
Rendezvous Hash	~K/N on change	Medium	When order matters
Jump Hash	~K/N on change	Low	Memory-constrained

## Key Takeaways

1. **Consistent hashing minimizes data movement** when scaling clusters
2. **Virtual nodes ensure even distribution** and support heterogeneous hardware
3. **The hash ring is a fundamental building block** for distributed databases and caches
4. **Replication integrates naturally** by using successor nodes
5. **Real systems like DynamoDB, Cassandra, and Redis** all use variants of consistent hashing



## Practical Exercise

**Scenario:** Design a distributed cache for a social media platform.

**Requirements:** - 10 cache servers initially - Must scale to 50 servers - Session data must be sticky (same user → same server) - Handle server failures gracefully

**Questions:** 1. How many virtual nodes per server? 2. How to handle user session stickiness? 3. What happens when a server fails? 4. How to migrate data when adding servers?

## Quick Reference

### Hash Ring Operations

```
# Add server  
ring.add_server("server-new")  
  
# Remove server  
ring.remove_server("server-old")  
  
# Find server for key  
server = ring.get_server("my-key")  
  
# Get replicas  
replicas = ring.get_servers_for_replication("my-key", n=3)
```

### Key Formulas

Keys remapped (modulo):  $(N-1)/N \approx 100\%$  for large N

Keys remapped (consistent):  $K/N$  where K=total keys, N=nodes

Virtual nodes memory: nodes × vnodes × ~100 bytes

Recommended vnodes: 100-256 per physical node

Session End - Thank You! 🎉

#consistent-hashing #distributed-systems #partitioning #architecture #session-notes