Distributed Transactions

Session Overview

ATTRIBUTE	DETAILS
Duration	60 minutes
Level	Intermediate to Advanced
Prerequisites	Database fundamentals, microservices basics

Agenda

| TIME	TOPIC
0-5 min	Introduction & ACID in distributed systems
5-20 min	Two-Phase Commit (2PC) & Three-Phase Commit (3PC)
20-35 min	Saga Pattern & Compensation
35-50 min	Eventual Consistency & Outbox Pattern
50-60 min	Practical Exercise & Discussion
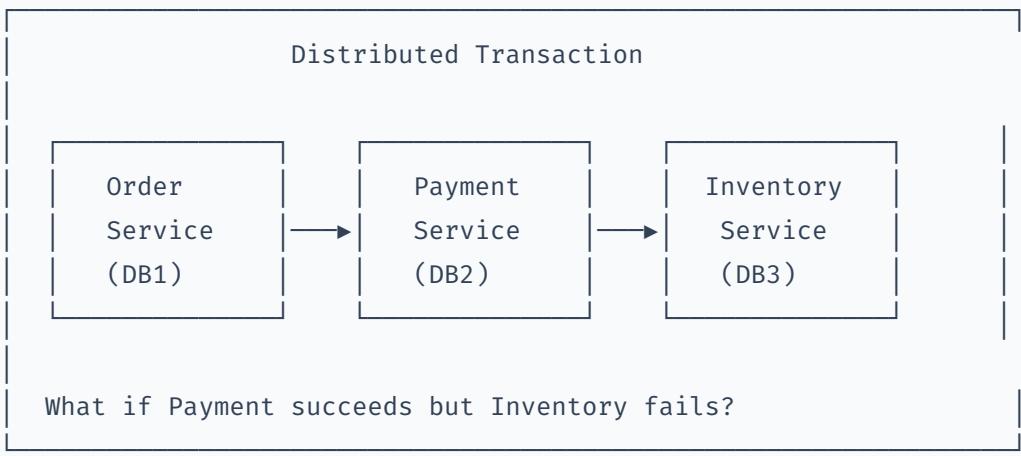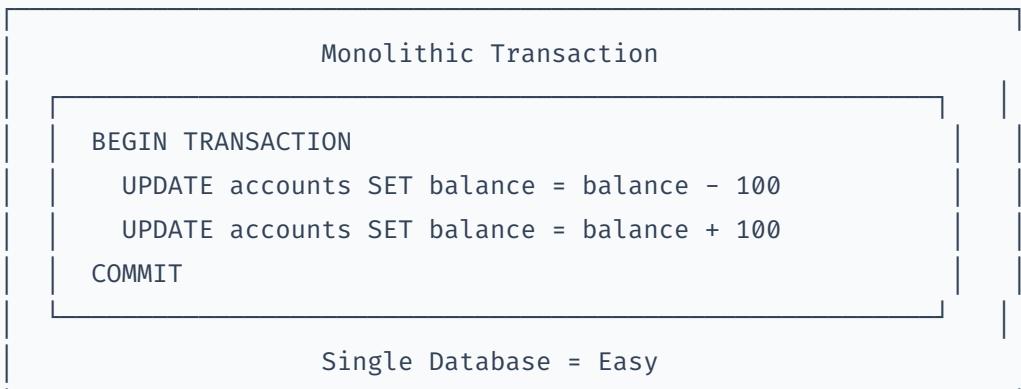
Learning Objectives

By the end of this session, you will be able to: - Understand the challenges of transactions across distributed systems - Compare 2PC, 3PC, and Saga patterns with their trade-offs - Design compensation strategies for failed transactions - Implement eventual consistency patterns in microservices - Choose the right transaction pattern for different use cases

# 1. The Distributed Transaction Challenge

Why Distributed Transactions Are Hard

```
┌──────────────────────────────────────────────────────────┐
│                  Monolithic Transaction                  │
│  ┌────────────────────────────────────────────────────┐  │
│  │  BEGIN TRANSACTION                                 │  │
│  │    UPDATE accounts SET balance = balance - 100    │  │
│  │    UPDATE accounts SET balance = balance + 100    │  │
│  │  COMMIT                                            │  │
│  └────────────────────────────────────────────────────┘  │
│                  Single Database = Easy                  │
└──────────────────────────────────────────────────────────┘


┌──────────────────────────────────────────────────────────┐
│                 Distributed Transaction                  │
│                                                          │
│  ┌─────────────┐      ┌─────────────┐      ┌─────────────┐  │
│  │   Order     │      │   Payment   │      │  Inventory  │  │
│  │   Service   │─────▶│   Service   │─────▶│   Service   │  │
│  │   (DB1)     │      │   (DB2)     │      │   (DB3)     │  │
│  └─────────────┘      └─────────────┘      └─────────────┘  │
│                                                          │
│  What if Payment succeeds but Inventory fails?          │
└──────────────────────────────────────────────────────────┘
```
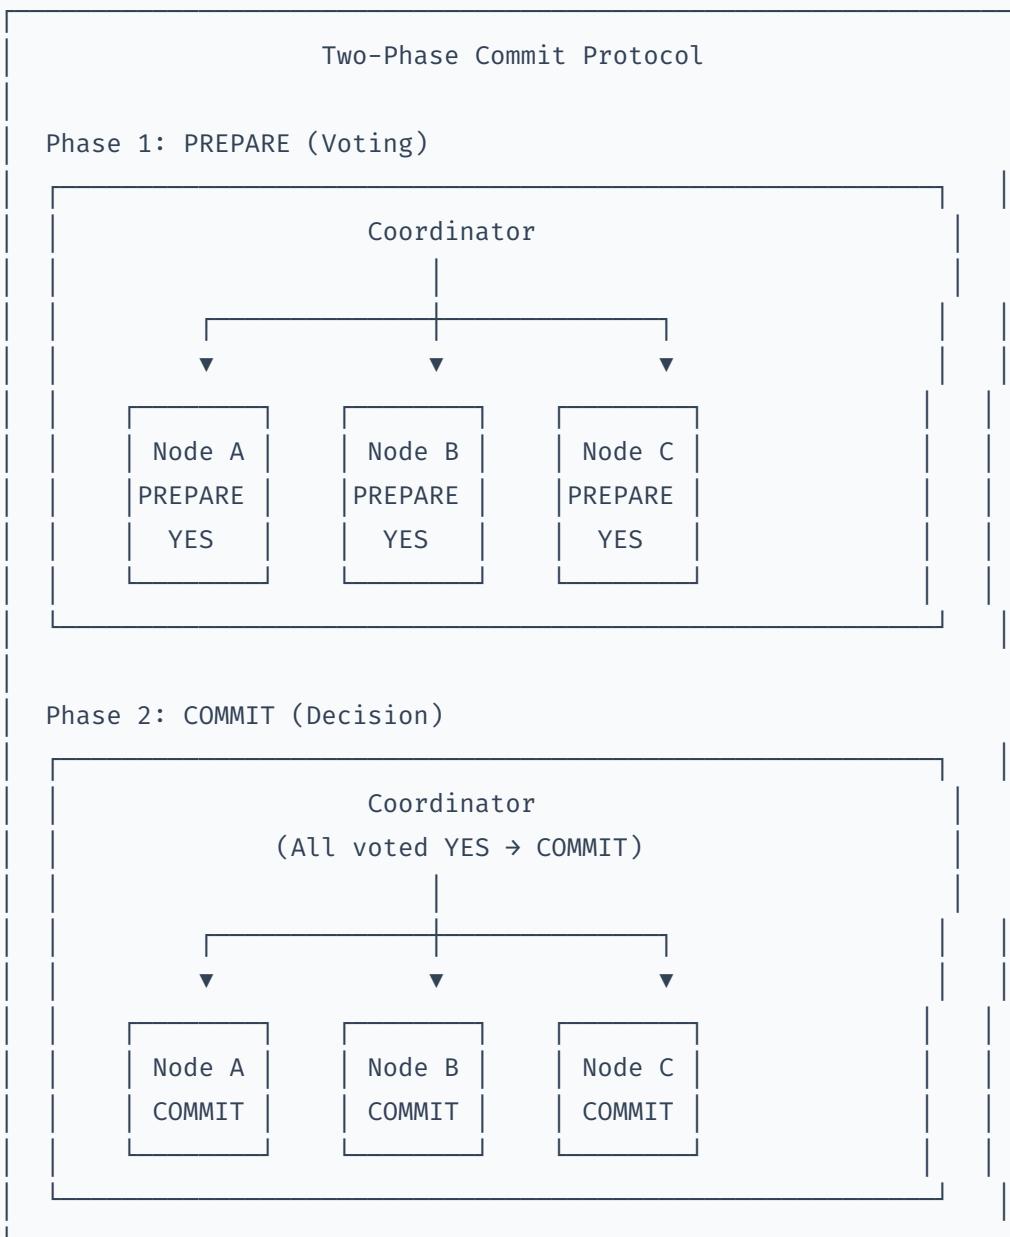
CAP Theorem Implications

PROPERTY	DESCRIPTION	IMPACT ON TRANSACTIONS
Consistency	All nodes see same data	Strong consistency = blocking
Availability	System always responds	May return stale data
Partition Tolerance	Works despite network splits	Must choose C or A

2. Two-Phase Commit (2PC)

How 2PC Works

```
┌──────────────────────────────────────────────────────────┐
│                Two-Phase Commit Protocol                   │
│                                                            │
│  Phase 1: PREPARE (Voting)                                 │
│  ┌──────────────────────────────────────────────────────┐ │
│  │                   Coordinator                          │ │
│  │                        │                               │ │
│  │          ┌─────────────┼─────────────┐                 │ │
│  │          ▼             ▼             ▼                 │ │
│  │    ┌──────────┐  ┌──────────┐  ┌──────────┐           │ │
│  │    │ Node A   │  │ Node B   │  │ Node C   │           │ │
│  │    │ PREPARE  │  │ PREPARE  │  │ PREPARE  │           │ │
│  │    │  YES     │  │  YES     │  │  YES     │           │ │
│  │    └──────────┘  └──────────┘  └──────────┘           │ │
│  └──────────────────────────────────────────────────────┘ │
│                                                            │
│  Phase 2: COMMIT (Decision)                                │
│  ┌──────────────────────────────────────────────────────┐ │
│  │                   Coordinator                          │ │
│  │             (All voted YES → COMMIT)                   │ │
│  │                        │                               │ │
│  │          ┌─────────────┼─────────────┐                 │ │
│  │          ▼             ▼             ▼                 │ │
│  │    ┌──────────┐  ┌──────────┐  ┌──────────┐           │ │
│  │    │ Node A   │  │ Node B   │  │ Node C   │           │ │
│  │    │ COMMIT   │  │ COMMIT   │  │ COMMIT   │           │ │
│  │    └──────────┘  └──────────┘  └──────────┘           │ │
│  └──────────────────────────────────────────────────────┘ │
└──────────────────────────────────────────────────────────┘
```

2PC Implementation Example

```java
public class TwoPhaseCommitCoordinator {
    private final List<TransactionParticipant> participants;
    private final TransactionLog transactionLog;

    public boolean executeTransaction(Transaction tx) {
        String txId = tx.getId();

        // Phase 1: Prepare
        transactionLog.logPrepare(txId);
        List<Boolean> votes = new ArrayList<>();

        for (TransactionParticipant participant : participants) {
            try {
                boolean vote = participant.prepare(tx);
                votes.add(vote);

                if (!vote) {
                    // Any NO vote → abort
                    return abort(txId);
                }
            } catch (Exception e) {
                // Timeout or failure → abort
                return abort(txId);
            }
        }

        // Phase 2: Commit (all voted YES)
        transactionLog.logCommit(txId);

        for (TransactionParticipant participant : participants) {
            try {
                participant.commit(txId);
            } catch (Exception e) {
                // Must retry until success (commit is durable)
                retryCommit(participant, txId);
            }
        }

        return true;
    }

    private boolean abort(String txId) {
```

```java
                transactionLog.logAbort(txId);

            for (TransactionParticipant participant : participants) {
                try {
                    participant.rollback(txId);
                } catch (Exception e) {
                    // Retry rollback
                    retryRollback(participant, txId);
                }
            }

            return false;
        }
    }

    public interface TransactionParticipant {
        boolean prepare(Transaction tx);  // Returns YES/NO vote
        void commit(String txId);
        void rollback(String txId);
    }
```
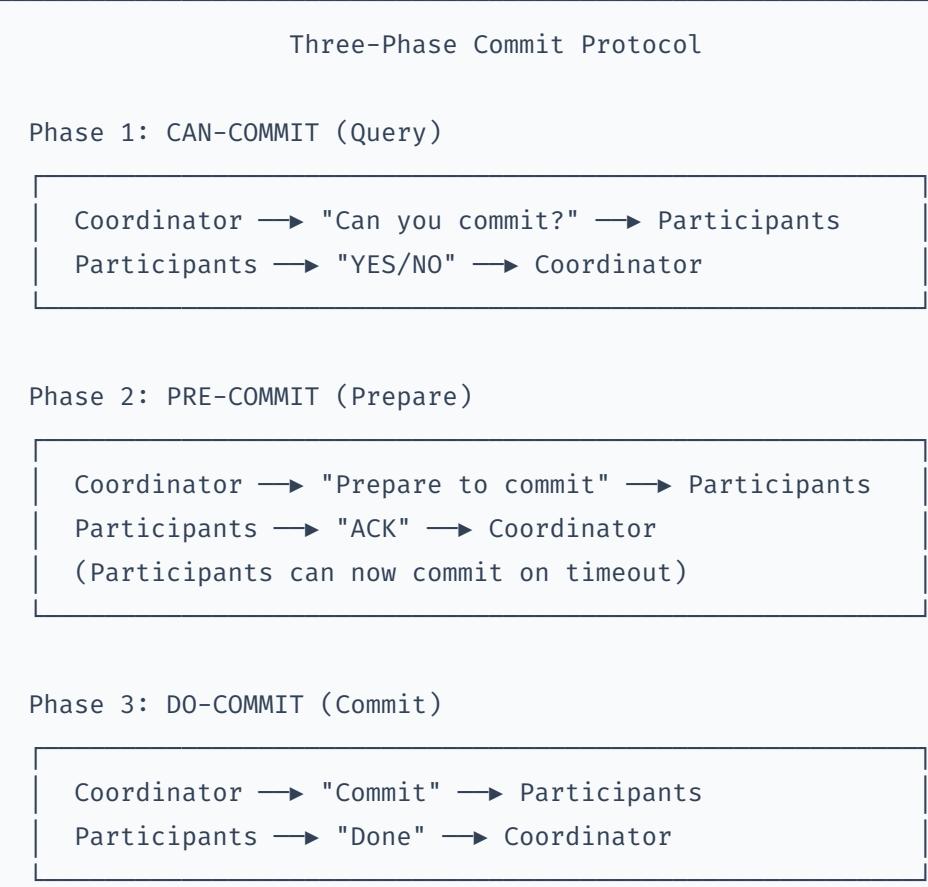
2PC Problems

PROBLEM	DESCRIPTION	IMPACT
Blocking	Participants hold locks during voting	Reduced throughput
Coordinator Failure	Single point of failure	Participants stuck in uncertain state
Network Partition	Can't reach consensus	Transaction hangs

# 3. Three-Phase Commit (3PC)

3PC Adds Pre-Commit Phase

```
┌─────────────────────────────────────────────────────────────┐
│                 Three-Phase Commit Protocol                  │
│                                                              │
│  Phase 1: CAN-COMMIT (Query)                                 │
│  ┌─────────────────────────────────────────────────────┐    │
│  │  Coordinator ──▶ "Can you commit?" ──▶ Participants  │    │
│  │  Participants ──▶ "YES/NO" ──▶ Coordinator           │    │
│  └─────────────────────────────────────────────────────┘    │
│                                                              │
│  Phase 2: PRE-COMMIT (Prepare)                               │
│  ┌─────────────────────────────────────────────────────┐    │
│  │  Coordinator ──▶ "Prepare to commit" ──▶ Participants│    │
│  │  Participants ──▶ "ACK" ──▶ Coordinator              │    │
│  │  (Participants can now commit on timeout)            │    │
│  └─────────────────────────────────────────────────────┘    │
│                                                              │
│  Phase 3: DO-COMMIT (Commit)                                 │
│  ┌─────────────────────────────────────────────────────┐    │
│  │  Coordinator ──▶ "Commit" ──▶ Participants           │    │
│  │  Participants ──▶ "Done" ──▶ Coordinator             │    │
│  └─────────────────────────────────────────────────────┘    │
└─────────────────────────────────────────────────────────────┘
```

2PC vs 3PC Comparison

| ASPECT	2PC	3PC
Phases	2	3
Blocking	Yes (indefinite)	Limited (timeout-based)
Coordinator Failure	Participants stuck	Participants can decide
Network Partitions	Problematic	Still problematic
Complexity	Lower	Higher
Latency	Lower	Higher
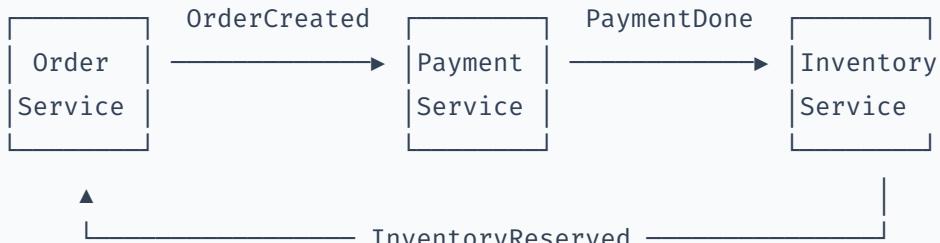
4. Saga Pattern

Saga: Sequence of Local Transactions

```
                        Saga Pattern

  Forward Flow (Happy Path):

  |     T1    |    |    T2    |    |    T3    |    |    T4    |
  |   Order   | -> |  Payment | -> | Inventory| -> | Shipping |


  Compensation Flow (T3 fails):

  |    C1    |    |    C2    |    |    T3    |   ✗
  |  Cancel  | <- |  Refund  | <- |  Failed  |
  |   Order  |    |  Payment |    |          |
```

Choreography vs Orchestration
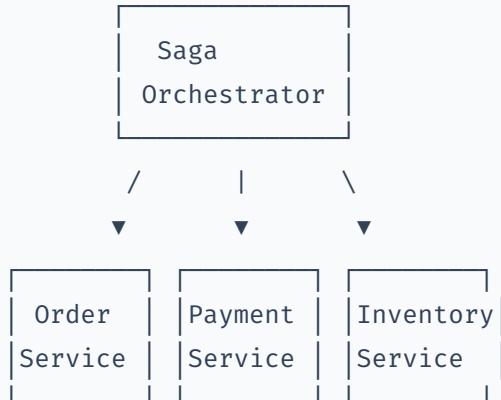
```
+-----------------------------------------------------------------+
|                  Choreography (Event-Driven)                    |
|                                                                 |
|                     OrderCreated            PaymentDone         |
|   +---------+                  +---------+              +---------+ |
|   | Order   | ---------------> |Payment  | -----------> |Inventory| |
|   |Service  |                  |Service  |              |Service  | |
|   +---------+                  +---------+              +---------+ |
|        ^                                                    |      |
|        +----------------- InventoryReserved ----------------+      |
|                                                                 |
|   Pros: Loose coupling, no single point of failure              |
|   Cons: Hard to track, complex failure handling                 |
+-----------------------------------------------------------------+


+-----------------------------------------------------------------+
|                  Orchestration (Central Control)                |
|                                                                 |
|                     +-------------+                             |
|                     |    Saga     |                             |
|                     |Orchestrator |                             |
|                     +-------------+                             |
|                       /    |    \                               |
|                      v     v     v                              |
|                 +-------+ +-------+ +---------+                  |
|                 | Order | |Payment| |Inventory|                 |
|                 |Service| |Service| |Service  |                 |
|                 +-------+ +-------+ +---------+                  |
|                                                                 |
|   Pros: Easy to track, centralized logic                        |
|   Cons: Single point of failure, tighter coupling               |
+-----------------------------------------------------------------+
```

Saga Orchestrator Implementation

```java
public class OrderSagaOrchestrator {
    private final OrderService orderService;
    private final PaymentService paymentService;
    private final InventoryService inventoryService;
    private final ShippingService shippingService;
    private final SagaStateStore stateStore;

    public SagaResult executeSaga(OrderRequest request) {
        String sagaId = UUID.randomUUID().toString();
        SagaState state = new SagaState(sagaId, SagaStep.ORDER_CREATED);

        try {
            // Step 1: Create Order
            Order order = orderService.createOrder(request);
            state.setOrderId(order.getId());
            state.setStep(SagaStep.ORDER_CREATED);
            stateStore.save(state);

            // Step 2: Process Payment
            PaymentResult payment = paymentService.processPayment(
                order.getId(), request.getPaymentInfo()
            );
            state.setPaymentId(payment.getId());
            state.setStep(SagaStep.PAYMENT_PROCESSED);
            stateStore.save(state);

            // Step 3: Reserve Inventory
            InventoryReservation reservation = inventoryService.reserve(
                order.getItems()
            );
            state.setReservationId(reservation.getId());
            state.setStep(SagaStep.INVENTORY_RESERVED);
            stateStore.save(state);

            // Step 4: Create Shipment
            Shipment shipment = shippingService.createShipment(order);
            state.setShipmentId(shipment.getId());
            state.setStep(SagaStep.COMPLETED);
            stateStore.save(state);

            return SagaResult.success(sagaId);
```

```java
            } catch (Exception e) {
                return compensate(state, e);
            }
        }

        private SagaResult compensate(SagaState state, Exception cause) {
            // Compensate in reverse order
            switch (state.getStep()) {
                case INVENTORY_RESERVED:

inventoryService.releaseReservation(state.getReservationId());
                    // Fall through
                case PAYMENT_PROCESSED:
                    paymentService.refund(state.getPaymentId());
                    // Fall through
                case ORDER_CREATED:
                    orderService.cancelOrder(state.getOrderId());
                    break;
            }

            state.setStep(SagaStep.COMPENSATED);
            stateStore.save(state);

            return SagaResult.failed(state.getSagaId(), cause.getMessage());
        }
    }

    enum SagaStep {
        STARTED,
        ORDER_CREATED,
        PAYMENT_PROCESSED,
        INVENTORY_RESERVED,
        COMPLETED,
        COMPENSATED
    }
```

Choreography with Events

```java
// Order Service
@Service
public class OrderService {
    private final EventPublisher eventPublisher;

    @Transactional
    public Order createOrder(OrderRequest request) {
        Order order = orderRepository.save(new Order(request));

        // Publish event for next step
        eventPublisher.publish(new OrderCreatedEvent(
            order.getId(),
            order.getItems(),
            order.getTotalAmount()
        ));

        return order;
    }


    @EventListener
    public void onPaymentFailed(PaymentFailedEvent event) {
        // Compensation: cancel order
        Order order = orderRepository.findById(event.getOrderId());
        order.setStatus(OrderStatus.CANCELLED);
        orderRepository.save(order);
    }
}

// Payment Service
@Service
public class PaymentService {
    private final EventPublisher eventPublisher;

    @EventListener
    public void onOrderCreated(OrderCreatedEvent event) {
        try {
            Payment payment = processPayment(event);

            eventPublisher.publish(new PaymentCompletedEvent(
                event.getOrderId(),
                payment.getId()
            ));
```
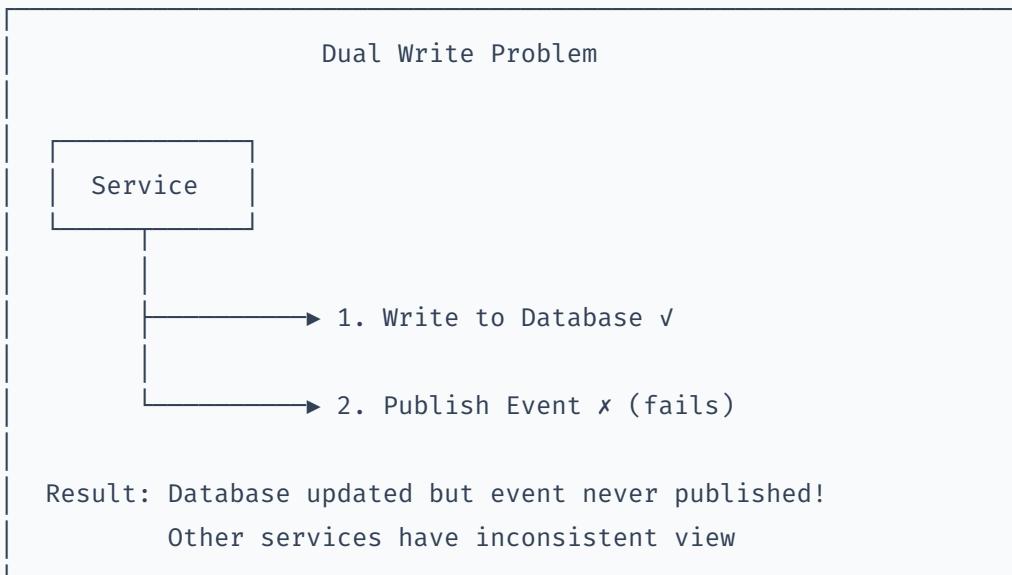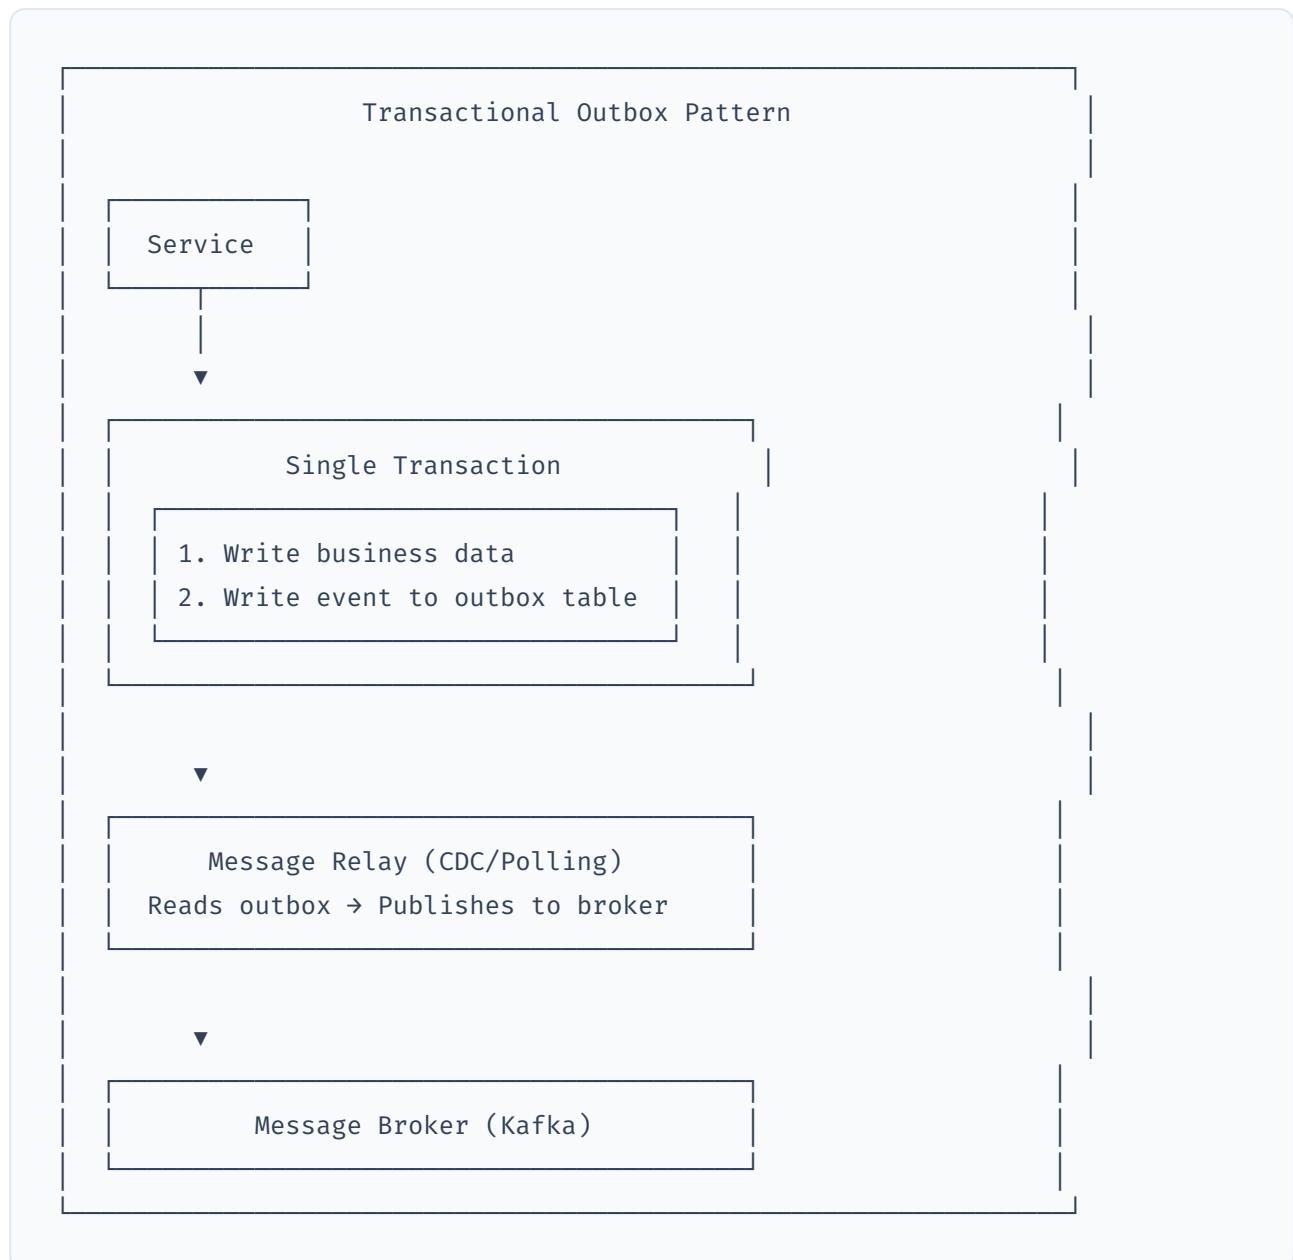
```
            } catch (PaymentException e) {
                eventPublisher.publish(new PaymentFailedEvent(
                    event.getOrderId(),
                    e.getMessage()
                ));
            }
        }
    }
```

5. Eventual Consistency & Outbox Pattern

The Dual Write Problem

```
┌─────────────────────────────────────────────────────┐
│                   Dual Write Problem                  │
│                                                       │
│   ┌───────────────┐                                   │
│   │   Service     │                                   │
│   └───────┬───────┘                                   │
│           │                                           │
│           ├──────────────▶ 1. Write to Database ✓     │
│           │                                           │
│           └──────────────▶ 2. Publish Event ✗ (fails) │
│                                                       │
│  Result: Database updated but event never published!  │
│        Other services have inconsistent view          │
└─────────────────────────────────────────────────────┘
```

Transactional Outbox Pattern

```
┌─────────────────────────────────────────────────────────────┐
│                 Transactional Outbox Pattern                  │
│                                                               │
│  ┌─────────────────┐                                          │
│  │  Service        │                                          │
│  └─────────────────┘                                          │
│          │                                                    │
│          ▼                                                    │
│  ┌──────────────────────────────────────────┐                │
│  │           Single Transaction        │     │                │
│  │   ┌─────────────────────────────┐   │     │                │
│  │   │ 1. Write business data      │   │     │                │
│  │   │ 2. Write event to outbox table │ │   │                │
│  │   └─────────────────────────────┘   │     │                │
│  └──────────────────────────────────────────┘                │
│                                                               │
│          ▼                                                    │
│  ┌──────────────────────────────────────────┐                │
│  │    Message Relay (CDC/Polling)           │                 │
│  │ Reads outbox → Publishes to broker       │                 │
│  └──────────────────────────────────────────┘                │
│                                                               │
│          ▼                                                    │
│  ┌──────────────────────────────────────────┐                │
│  │        Message Broker (Kafka)            │                 │
│  └──────────────────────────────────────────┘                │
└─────────────────────────────────────────────────────────────┘
```

Outbox Implementation

```java
@Entity
@Table(name = "outbox_events")
public class OutboxEvent {
    @Id
    private String id;

    private String aggregateType;  // e.g., "Order"
    private String aggregateId;    // e.g., order ID
    private String eventType;      // e.g., "OrderCreated"

    @Column(columnDefinition = "TEXT")
    private String payload;        // JSON event data

    private Instant createdAt;
    private boolean published;
}


@Service
public class OrderService {
    private final OrderRepository orderRepository;
    private final OutboxRepository outboxRepository;

    @Transactional  // Single transaction!
    public Order createOrder(OrderRequest request) {
        // 1. Save business entity
        Order order = new Order(request);
        orderRepository.save(order);

        // 2. Save event to outbox (same transaction)
        OutboxEvent event = new OutboxEvent();
        event.setId(UUID.randomUUID().toString());
        event.setAggregateType("Order");
        event.setAggregateId(order.getId());
        event.setEventType("OrderCreated");
        event.setPayload(toJson(new OrderCreatedEvent(order)));
        event.setCreatedAt(Instant.now());
        event.setPublished(false);

        outboxRepository.save(event);

        return order;
    }
```

```java
        }

        // Message Relay (runs separately)
        @Scheduled(fixedDelay = 1000)
        public void publishOutboxEvents() {
            List<OutboxEvent> events = outboxRepository
                .findByPublishedFalseOrderByCreatedAt();

            for (OutboxEvent event : events) {
                try {
                    kafkaTemplate.send(
                        event.getAggregateType(),
                        event.getAggregateId(),
                        event.getPayload()
                    );

                    event.setPublished(true);
                    outboxRepository.save(event);

                } catch (Exception e) {
                    // Will retry on next poll
                    log.error("Failed to publish event: {}", event.getId(), e);
                }
            }
        }
```

Change Data Capture (CDC) with Debezium

```
# Debezium connector configuration
{
  "name": "outbox-connector",
  "config": {
    "connector.class": "io.debezium.connector.postgresql.PostgresConnector",
    "database.hostname": "postgres",
    "database.port": "5432",
    "database.user": "debezium",
    "database.password": "secret",
    "database.dbname": "orders",
    "table.include.list": "public.outbox_events",
    "transforms": "outbox",
    "transforms.outbox.type": "io.debezium.transforms.outbox.EventRouter",
    "transforms.outbox.route.topic.replacement": "${routedByValue}",
    "transforms.outbox.table.field.event.key": "aggregate_id",
    "transforms.outbox.table.field.event.payload": "payload"
  }
}
```

6. Comparison of Approaches

ASPECT	2PC	SAGA	EVENTUAL CONSISTENCY
Consistency	Strong	Eventual	Eventual
Isolation	Full	None (dirty reads possible)	None
Latency	High (blocking)	Medium	Low
Complexity	Medium	High	Medium
Scalability	Low	High	High
Failure Handling	Automatic rollback	Compensation logic	Idempotent handlers
Use Case	Financial systems	E-commerce, booking	Social media, analytics

Key Takeaways

1. **2PC provides strong consistency** but blocks and doesn't scale well

2. **Sagas trade consistency for availability** using compensation

3. **Choreography is loosely coupled** but hard to track and debug

4. **Orchestration centralizes logic** but creates a single point of failure

5. **Outbox pattern solves dual-write** by using a single transaction

6. **CDC (Debezium) is more reliable** than polling for outbox relay

7. **Design for idempotency** - messages may be delivered multiple times

8. **Choose based on requirements** - not all systems need strong consistency

Practical Exercise

Design a distributed transaction strategy for an e-commerce order flow:

Requirements: - Order creation, payment processing, inventory reservation, shipping - Payment failures should cancel the order - Inventory failures should refund payment and cancel order - System should handle partial failures gracefully

Tasks: 1. Draw the saga flow with compensation steps 2. Decide between choreography and orchestration (justify your choice) 3. Implement the outbox pattern for reliable event publishing 4. Handle idempotency for payment processing

Discussion Points: - What happens if compensation fails? - How do you handle long-running transactions (e.g., shipping takes days)? - How do you provide transaction status to users?