

Rate Limiter - Discussion Topics

Architecture & Design

1. When would you choose token bucket vs sliding window log?
 - Memory efficiency considerations
 - Burst handling requirements
 - Precision vs performance trade-offs
2. How do you implement rate limiting in a distributed system?
 - Centralized vs distributed counters
 - Redis-based implementations
 - Eventual consistency challenges
3. Where should rate limiting be implemented in the architecture?
 - API Gateway level
 - Application level
 - Database level
 - Trade-offs of each approach

Real-World Scenarios

4. Design a rate limiter for a public API with multiple pricing tiers
 - Different limits per tier
 - Graceful degradation
 - Quota management
5. How would you handle rate limiting for a viral event?
 - Sudden traffic spikes
 - Fair distribution of resources
 - Protecting critical services
6. Design rate limiting for a payment processing system
 - Fraud prevention
 - User experience balance
 - Compliance requirements

Performance & Optimization

7. How do you minimize latency impact of rate limiting?
 - Local caching strategies
 - Async counter updates
 - Optimistic vs pessimistic approaches
8. What metrics would you monitor for rate limiting effectiveness?
 - Rejection rates
 - Latency percentiles
 - False positive rates

Advanced Topics

9. How do you handle rate limiting across multiple regions?
 - Global vs regional limits
 - Synchronization strategies
 - Consistency models
10. What's your approach to adaptive rate limiting?
 - Dynamic limit adjustment
 - Load-based throttling
 - Machine learning approaches
11. How do you prevent rate limit bypass attempts?

- IP rotation detection
- Fingerprinting techniques
- Behavioral analysis