

Service Discovery - Discussion Topics

Architecture & Design

1. When to use client-side vs server-side service discovery?
 - Trade-offs between complexity and flexibility
 - Impact on client libraries and coupling
 - Load balancing considerations
2. How do you choose between Consul, etcd, and ZooKeeper?
 - Consistency models (CP vs AP)
 - Operational complexity
 - Feature sets and ecosystem support
3. What are the risks of DNS-based service discovery?
 - TTL and caching issues
 - Propagation delays
 - Load balancing limitations

Real-World Scenarios

4. Design service discovery for a multi-region deployment
 - Regional vs global registries
 - Failover strategies
 - Latency-aware routing
5. How would you handle a service registry failure?
 - Graceful degradation
 - Local caching strategies
 - Recovery procedures

Health Checking & Reliability

6. What's your approach to health check design?
 - Shallow vs deep health checks
 - Check frequency and timeout tuning
 - Cascading failure prevention
7. How do you prevent flapping services?
 - Debouncing strategies
 - Gradual rollout/rollback
 - Circuit breaker integration

Advanced Topics

8. Service mesh vs traditional service discovery: When to use each?
 - Sidecar proxy trade-offs
 - Operational complexity
 - Feature requirements
9. How do you handle service discovery during deployments?
 - Blue-green deployment integration
 - Canary release support
 - Connection draining strategies