

Service Discovery

Session Overview

Duration	Level	Prerequisites
60 min	Intermediate	Basic microservices concepts, networking fundamentals

Learning Objectives

By the end of this session, you will be able to:

- Understand the need for service discovery in distributed systems
- Compare client-side vs server-side discovery patterns
- Evaluate different service registry technologies (Consul, etcd, ZooKeeper)
- Implement health checking and failure detection
- Design DNS-based service discovery solutions

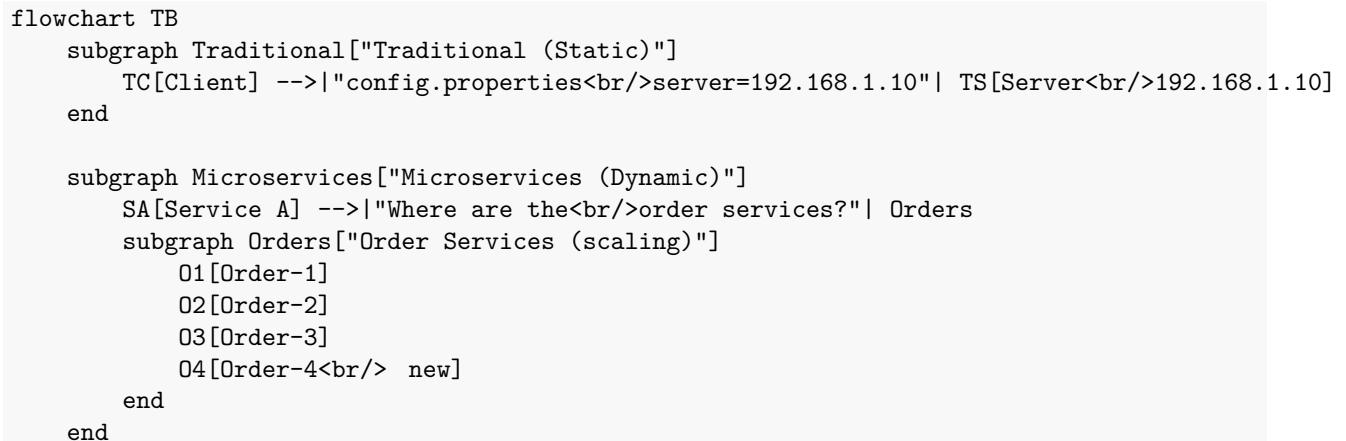
Agenda

Time	Topic
0-10 min	Why Service Discovery?
10-25 min	Discovery Patterns
25-40 min	Service Registry Technologies
40-50 min	Health Checking & DNS-Based Discovery
50-60 min	Practical Exercise & Q&A

1. Why Service Discovery?

The Problem

In traditional monolithic applications, service locations are typically static and configured at deployment time. In microservices architectures, this approach breaks down:



Challenges in Dynamic Environments

1. **Auto-scaling:** Instances are added/removed based on load
2. **Failures:** Instances crash and restart on different hosts
3. **Deployments:** Rolling updates change instance locations
4. **Containerization:** Containers get dynamic IPs

Service Discovery Components

```
flowchart TB
    subgraph SD ["Service Discovery System"]
        SR["Service Registry<br/>• Register<br/>• Deregister<br/>• Store"]
        DC["Service Discovery Client<br/>• Query<br/>• Subscribe<br/>• Cache"]
        HC["Health Checking<br/>• Liveness<br/>• Readiness<br/>• Monitor"]
    end
```

2. Discovery Patterns

Client-Side Discovery

The client is responsible for determining the network locations of available service instances and load balancing requests across them.

```
flowchart LR
    subgraph Client ["Client"]
        DC["Discovery<br/>Client"]
        LB["Load<br/>Balancer"]
    end

    subgraph Registry ["Service Registry"]
        SR["order-svc:<br/>• 10.0.1.5<br/>• 10.0.1.6<br/>• 10.0.1.7"]
    end

    SI[Service Instance]

    DC -->|"1. Query"| SR
    SR -->|"2. Return instances"| DC
    DC --> LB
    LB -->|"3. Direct request"| SI
```

Implementation Example (Java with Netflix Eureka):

```
@Configuration
@EnableDiscoveryClient
public class ServiceDiscoveryConfig {

    @Bean
    @LoadBalanced
    public RestTemplate restTemplate() {
        return new RestTemplate();
    }

    @Service
    public class OrderServiceClient {

        @Autowired
        private RestTemplate restTemplate;

        @Autowired
        private DiscoveryClient discoveryClient;
    }
}
```

```

// Using service name instead of hardcoded URL
public Order getOrder(String orderId) {
    return restTemplate.getForObject(
        "http://order-service/orders/" + orderId,
        Order.class
    );
}

// Manual discovery for more control
public List<ServiceInstance> getOrderServiceInstances() {
    return discoveryClient.getInstances("order-service");
}
}

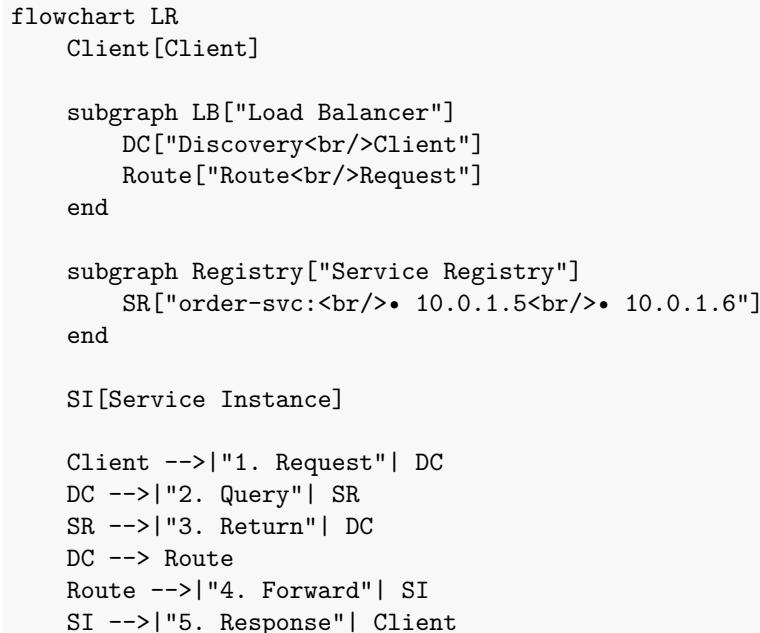
```

Pros: - Client has full control over load balancing strategy - No single point of failure in the load balancer
- Can implement sophisticated routing (e.g., zone-aware)

Cons: - Discovery logic coupled with client code - Must implement for each programming language - Clients must handle registry unavailability

Server-Side Discovery

A load balancer or router handles service discovery, and clients make requests to the load balancer.



Implementation Example (Kubernetes Service):

```

# Kubernetes Service acts as server-side discovery
apiVersion: v1
kind: Service
metadata:
  name: order-service
spec:
  selector:
    app: order

```

```

ports:
  - protocol: TCP
    port: 80
    targetPort: 8080
  type: ClusterIP

---

# Client just uses the service name
apiVersion: apps/v1
kind: Deployment
metadata:
  name: api-gateway
spec:
  template:
    spec:
      containers:
        - name: gateway
          env:
            - name: ORDER_SERVICE_URL
              value: "http://order-service" # DNS resolves to service

```

Pros: - Simpler client implementation - Language-agnostic - Centralized load balancing policies

Cons: - Load balancer can become a bottleneck - Additional network hop - Must ensure load balancer high availability

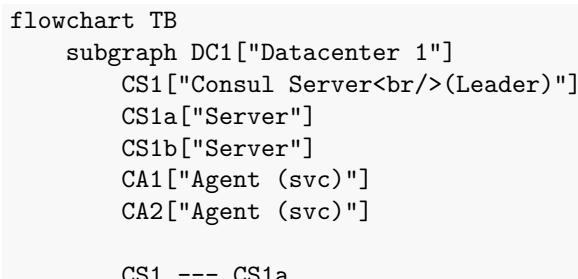
Pattern Comparison

Aspect	Client-Side	Server-Side
Client Complexity	High	Low
Network Hops	Fewer	More
Load Balancer	Distributed	Centralized
Language Support	Per-language	Universal
Flexibility	High	Medium
Examples	Netflix Ribbon, gRPC	Kubernetes, AWS ALB

3. Service Registry Technologies

Consul

HashiCorp Consul provides service discovery with built-in health checking, KV store, and multi-datacenter support.



```

    CS1 --- CS1b
    CS1a --- CA1
    CS1b --- CA2
end

subgraph DC2["Datacenter 2"]
    CS2["Consul Server<br/>(Leader)"]
    CS2a["Server"]
    CS2b["Server"]
    CA3["Agent (svc)"]
    CA4["Agent (svc)"]

    CS2 --- CS2a
    CS2 --- CS2b
    CS2a --- CA3
    CS2b --- CA4
end

CS1 <-->| "WAN Gossip" | CS2

```

Service Registration:

```

// service.json
{
  "service": {
    "name": "order-service",
    "id": "order-service-1",
    "port": 8080,
    "tags": ["v1", "primary"],
    "check": {
      "http": "http://localhost:8080/health",
      "interval": "10s",
      "timeout": "5s"
    }
  }
}

# Register service
consul services register service.json

# Query service
consul catalog services
consul catalog nodes -service=order-service

# DNS query
dig @127.0.0.1 -p 8600 order-service.service.consul

```

Go Client Example:

```

package main

import (
    "fmt"
    "github.com/hashicorp/consul/api"
)

```

```

func main() {
    // Create client
    config := api.DefaultConfig()
    client, _ := api.NewClient(config)

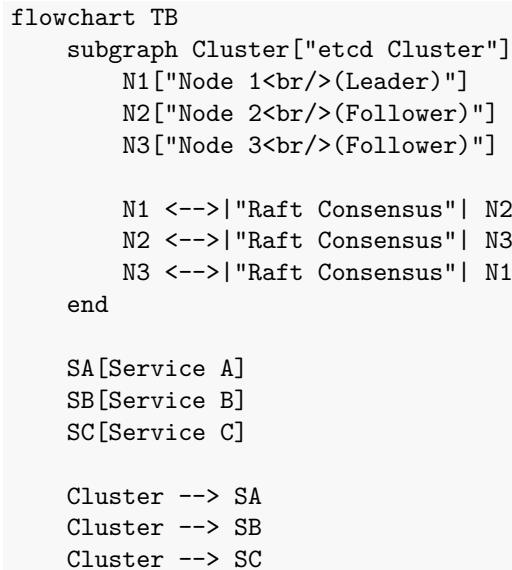
    // Register service
    registration := &api.AgentServiceRegistration{
        ID:      "order-service-1",
        Name:    "order-service",
        Port:    8080,
        Check:  &api.AgentServiceCheck{
            HTTP:    "http://localhost:8080/health",
            Interval: "10s",
        },
    }
    client.Agent().ServiceRegister(registration)

    // Discover services
    services, _, _ := client.Health().Service("order-service", "", true, nil)
    for _, svc := range services {
        fmt.Printf("Found: %s:%d\n", svc.Service.Address, svc.Service.Port)
    }
}

```

etcd

etcd is a distributed key-value store that provides a reliable way to store data across a cluster. It's the backbone of Kubernetes service discovery.



Service Registration with etcd:

```

import etcd3
import json
import time
from threading import Thread

```

```

class ServiceRegistry:
    def __init__(self, etcd_host='localhost', etcd_port=2379):
        self.client = etcd3.client(host=etcd_host, port=etcd_port)
        self.lease = None

    def register(self, service_name, instance_id, host, port, ttl=30):
        """Register service with TTL-based lease"""
        key = f"/services/{service_name}/{instance_id}"
        value = json.dumps({
            "host": host,
            "port": port,
            "registered_at": time.time()
        })

        # Create lease for automatic deregistration
        self.lease = self.client.lease(ttl)
        self.client.put(key, value, lease=self.lease)

        # Start heartbeat thread
        Thread(target=self._heartbeat, daemon=True).start()

    def _heartbeat(self):
        """Keep lease alive"""
        while True:
            self.lease.refresh()
            time.sleep(10)

    def discover(self, service_name):
        """Get all instances of a service"""
        prefix = f"/services/{service_name}/"
        instances = []

        for value, metadata in self.client.get_prefix(prefix):
            instances.append(json.loads(value))

        return instances

    def watch(self, service_name, callback):
        """Watch for service changes"""
        prefix = f"/services/{service_name}/"
        events_iterator, cancel = self.client.watch_prefix(prefix)

        for event in events_iterator:
            callback(event)

# Usage
registry = ServiceRegistry()
registry.register("order-service", "instance-1", "10.0.1.5", 8080)
instances = registry.discover("order-service")

```

ZooKeeper

Apache ZooKeeper provides distributed coordination and is used by many systems for service discovery.

```

flowchart TB
    subgraph ZK["ZooKeeper Data Model"]
        Root["/"]
        Services【"/services"】
        Config【"/config"】

        OrderSvc【"/order-service"】
        UserSvc【"/user-service"】
        DB【"/database"】

        01【"instance-001 (ephemeral)  
host:10.0.1.5, port:8080"】
        02【"instance-002 (ephemeral)  
host:10.0.1.6, port:8080"】
        03【"instance-003 (ephemeral)  
host:10.0.1.7, port:8080"】

        U1【"instance-001 (ephemeral)"】
        U2【"instance-002 (ephemeral)"】

        DBConfig【"url: jdbc:mysql://..."】

        Root --> Services
        Root --> Config
        Services --> OrderSvc
        Services --> UserSvc
        Config --> DB
        OrderSvc --> 01
        OrderSvc --> 02
        OrderSvc --> 03
        UserSvc --> U1
        UserSvc --> U2
        DB --> DBConfig
    end

```

Note: Ephemeral nodes are automatically deleted when the session ends (service crashes)

Java Implementation with Curator:

```

import org.apache.curator.framework.CuratorFramework;
import org.apache.curator.framework.CuratorFrameworkFactory;
import org.apache.curator.retry.ExponentialBackoffRetry;
import org.apache.curator.x.discovery.*;

public class ZookeeperServiceDiscovery {

    private final ServiceDiscovery<ServiceInstance> serviceDiscovery;
    private final CuratorFramework client;

    public ZookeeperServiceDiscovery(String zkConnection) throws Exception {
        client = CuratorFrameworkFactory.newClient(
            zkConnection,
            new ExponentialBackoffRetry(1000, 3)
        );
        client.start();

        serviceDiscovery = ServiceDiscoveryBuilder.builder(ServiceInstance.class)
            .client(client)
    }
}

```

```

        .basePath("/services")
        .build();
    serviceDiscovery.start();
}

public void register(String serviceName, String host, int port) throws Exception {
    ServiceInstance<ServiceInstance> instance = ServiceInstance.<ServiceInstance>builder()
        .name(serviceName)
        .address(host)
        .port(port)
        .build();

    serviceDiscovery.registerService(instance);
}

public Collection<ServiceInstance<ServiceInstance>> discover(String serviceName)
    throws Exception {
    return serviceDiscovery.queryForInstances(serviceName);
}

// Watch for changes
public void watchService(String serviceName, ServiceCacheListener listener)
    throws Exception {
    ServiceCache<ServiceInstance> cache = serviceDiscovery
        .serviceCacheBuilder()
        .name(serviceName)
        .build();

    cache.addListener(listener);
    cache.start();
}
}

```

Technology Comparison

Feature	Consul	etcd	ZooKeeper
Consensus	Raft	Raft	ZAB
Health Checking	Built-in	External	External
Multi-DC	Native	Manual	Manual
DNS Interface	Yes	No	No
KV Store	Yes	Yes	Yes
ACL	Yes	Yes	Yes
Language	Go	Go	Java
Kubernetes	Supported	Native	Supported

4. Health Checking & DNS-Based Discovery

Health Check Types

```

flowchart LR
    subgraph Liveness["1. Liveness Check"]

```

```

    L_Q["Is the process running?"]
    LC[Checker] -->|"GET /health/live" | LS[Service]
    LS -->|"200 OK" | LC
end

subgraph Readiness["2. Readiness Check"]
    R_Q["Can it handle traffic?"]
    RC[Checker] -->|"GET /health/ready" | RS[Service]
    RS -->|"200 OK / 503" | RC
end

subgraph Startup["3. Startup Check"]
    S_Q["Has it finished initializing?"]
    SC[Checker] -->|"GET /health/startup" | SS[Service]
    SS -->|"200 OK / 503" | SC
end

```

Health Check Implementation:

```

@RestController
@RequestMapping("/health")
public class HealthController {

    @Autowired
    private DataSource dataSource;

    @Autowired
    private RedisTemplate<String, String> redisTemplate;

    @GetMapping("/live")
    public ResponseEntity<Map<String, String>> liveness() {
        // Simple check - is the process running?
        return ResponseEntity.ok(Map.of("status", "UP"));
    }

    @GetMapping("/ready")
    public ResponseEntity<Map<String, Object>> readiness() {
        Map<String, Object> health = new HashMap<>();
        boolean ready = true;

        // Check database
        try {
            dataSource.getConnection().isValid(5);
            health.put("database", "UP");
        } catch (Exception e) {
            health.put("database", "DOWN");
            ready = false;
        }

        // Check Redis
        try {
            redisTemplate.getConnectionFactory().getConnection().ping();
            health.put("redis", "UP");
        } catch (Exception e) {
    
```

```

        health.put("redis", "DOWN");
        ready = false;
    }

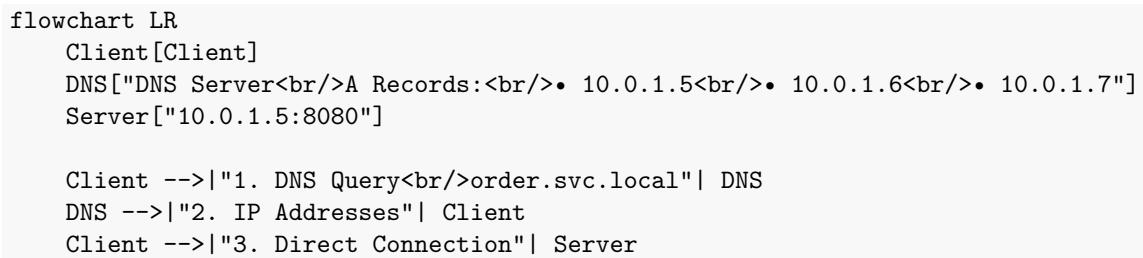
    health.put("status", ready ? "UP" : "DOWN");

    return ready
        ? ResponseEntity.ok(health)
        : ResponseEntity.status(503).body(health);
    }
}

```

DNS-Based Service Discovery

DNS-based discovery uses standard DNS protocols to resolve service names to IP addresses.



SRV Records (with port info):

```

_http._tcp.order.svc.local. 86400 IN SRV 0 5 8080 order1
_http._tcp.order.svc.local. 86400 IN SRV 0 5 8080 order2

```

Kubernetes DNS Discovery:

```

# Service creates DNS entries automatically
apiVersion: v1
kind: Service
metadata:
  name: order-service
  namespace: production
spec:
  selector:
    app: order
  ports:
    - port: 80
      targetPort: 8080

# DNS names created:
# - order-service.production.svc.cluster.local (full)
# - order-service.production (within cluster)
# - order-service (within same namespace)

// Go client using DNS
package main

import (
    "net"
    "net/http"
)

```

```

)

func main() {
    // Simple DNS resolution
    ips, _ := net.LookupHost("order-service.production.svc.cluster.local")

    // SRV record lookup for port info
    _, addrs, _ := net.LookupSRV("http", "tcp", "order-service.production.svc.cluster.local")

    for _, addr := range addrs {
        fmt.Printf("Target: %s, Port: %d\n", addr.Target, addr.Port)
    }

    // HTTP client with DNS
    client := &http.Client{}
    resp, _ := client.Get("http://order-service/api/orders")
}

```

5. Implementation Patterns

Self-Registration Pattern

Services register themselves with the registry on startup.

```

sequenceDiagram
    participant S as Service Instance
    participant R as Registry

    Note over S,R: Service Startup
    S->>R: 1. Register
    loop Periodic
        S->>R: 2. Heartbeat
    end

    Note over S,R: Service Shutdown
    S->>R: 3. Deregister

```

```

@Component
public class ServiceRegistration implements ApplicationListener<ApplicationReadyEvent> {

    @Autowired
    private ConsulClient consulClient;

    @Value("${spring.application.name}")
    private String serviceName;

    @Value("${server.port}")
    private int port;

    private String instanceId;

    @Override
    public void onApplicationEvent(ApplicationReadyEvent event) {

```

```

instanceId = serviceName + "-" + UUID.randomUUID().toString();

NewService service = new NewService();
service.setId(instanceId);
service.setName(serviceName);
service.setPort(port);
service.setCheck(createHealthCheck());

consulClient.agentServiceRegister(service);
}

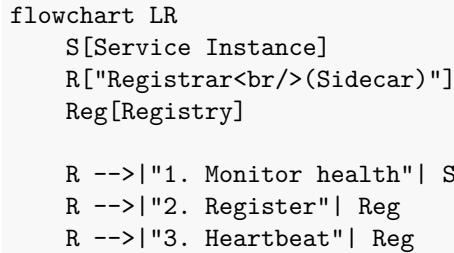
@PreDestroy
public void deregister() {
    consulClient.agentServiceDeregister(instanceId);
}

private NewService.Check createHealthCheck() {
    NewService.Check check = new NewService.Check();
    check.setHttp("http://localhost:" + port + "/health");
    check.setInterval("10s");
    return check;
}
}

```

Third-Party Registration Pattern

A separate registrar service handles registration on behalf of services.



Examples: Kubernetes, Netflix Prana, Registrar

Key Takeaways

1. **Service discovery is essential** for dynamic microservices environments where instances come and go
2. **Two main patterns:** Client-side (more control, more complexity) vs Server-side (simpler clients, centralized)
3. **Choose the right registry:** Consul for multi-DC, etcd for Kubernetes, ZooKeeper for existing Hadoop ecosystems
4. **Health checking is critical:** Implement liveness, readiness, and startup probes appropriately
5. **DNS-based discovery** provides simplicity and language-agnostic access but has caching challenges
6. **Consider failure modes:** What happens when the registry is unavailable? Cache locally and fail gracefully

Practical Exercise

Design Challenge: Multi-Region Service Discovery

Design a service discovery solution for a global e-commerce platform with the following requirements:

1. **Services deployed across 3 regions:** US-East, EU-West, Asia-Pacific
2. **Latency-sensitive routing:** Prefer local region, fallback to others
3. **Graceful degradation:** Continue operating if one region's registry fails
4. **Health-aware routing:** Only route to healthy instances

Consider: - Which discovery pattern (client-side vs server-side)? - Which registry technology? - How to handle cross-region discovery? - Caching strategy for registry data? - Failover behavior?

Deliverables: 1. Architecture diagram showing discovery flow 2. Sequence diagram for service registration 3. Failure scenario analysis

Q&A

Common questions to explore: - How does Kubernetes handle service discovery internally? - What's the difference between service mesh discovery and traditional discovery? - How do you migrate from static configuration to dynamic discovery? - What are the security considerations for service registries?