

Social Media System Design - Progressive Stages

Stage 2: Search with Elasticsearch

Architecture Overview

Added full-text search capability with Elasticsearch, migrated to NoSQL, and implemented event-driven indexing.

```
graph TB
    subgraph "Stage 2: Search-Enabled Architecture"
        Client[Client/Browser]
        API[API Server<br/>User CRUD<br/>Post CRUD<br/>Search API]

        subgraph "Data Layer"
            MongoDB[(NoSQL Database<br/>MongoDB)]
            ES[(Elasticsearch<br/>Search Index)]
        end

        subgraph "Event Processing"
            Kafka[Apache Kafka<br/>Event Stream]
            Indexer[Indexer Service<br/>Consumes Events]
            Cron[Cron Job<br/>Reindex Service]
        end

        Client -->|HTTP Requests| API
        API -->|Write Operations| MongoDB
        API -->|Search Queries| ES
        API -->|Publish Events| Kafka
        Kafka -->|Subscribe| Indexer
        Indexer -->|Update Index| ES
        MongoDB -->|Read Data| Cron
        Cron -->|Bulk Reindex| ES
    end

    style API fill:#4A90E2
    style MongoDB fill:#47A248
    style ES fill:#FEC514
    style Kafka fill:#231F20
    style Indexer fill:#FF6B6B
    style Cron fill:#9B59B6
```

Data Flow - Create Post with Indexing

```
sequenceDiagram
    participant C as Client
    participant A as API Server
    participant M as MongoDB
    participant K as Kafka
    participant I as Indexer Service
    participant E as Elasticsearch

    Note over C,E: Create Post Flow
    C->>A: POST /api/posts
    A->>M: Insert Post Document
```

```

M-->>A: Success (post_id)

par Async Indexing
  A->>K: Publish "post.created" Event
  K->>I: Consume Event
  I->>E: Index Post Document
  E-->>I: Indexed
end

A-->>C: 201 Created

```

Data Flow - Search Posts

```

sequenceDiagram
    participant C as Client
    participant A as API Server
    participant E as Elasticsearch
    participant M as MongoDB

    Note over C,M: Search Flow
    C->>A: GET /api/posts/search?q=keyword
    A->>E: Query: match(content, keyword)
    E->>E: Rank by relevance,<br/>recency, votes
    E-->>A: Post IDs + Scores

    opt Fetch Full Details
        A->>M: Get Posts by IDs
        M-->>A: Post Documents
    end

    A-->>C: 200 OK + Search Results

```

Reindexing Flow

```

sequenceDiagram
    participant Cron as Cron Job
    participant M as MongoDB
    participant E as Elasticsearch

    Note over Cron,E: Scheduled Reindex (Daily 2 AM)
    Cron->>Cron: Trigger Reindex Job
    Cron->>M: Fetch All Posts (Batch)
    M-->>Cron: Posts Data

    loop For Each Batch
        Cron->>E: Bulk Index Documents
        E-->>Cron: Batch Indexed
    end

    Cron->>Cron: Log Completion

```

Kafka Topics Structure

```
graph LR
    subgraph "Kafka Topics"
        T1[post.created]
        T2[post.updated]
        T3[post.deleted]
        T4[post.voted]
    end

    API[API Server] -->|Produce| T1
    API -->|Produce| T2
    API -->|Produce| T3
    API -->|Produce| T4

    T1 -->|Consume| Indexer[Indexer Service]
    T2 -->|Consume| Indexer
    T3 -->|Consume| Indexer
    T4 -->|Consume| Indexer

    style T1 fill:#2ECC71
    style T2 fill:#3498DB
    style T3 fill:#E74C3C
    style T4 fill:#F39C12
```

Stage 3: Load Balancing & Caching

Architecture Overview

Production-ready architecture with load balancing, caching layer, and horizontal scaling.

```
graph TB
    subgraph "Stage 3: Scaled Architecture with LB & Cache"
        Client[Clients/Browsers]
        LB[Load Balancer<br/>Nginx/AWS ALB]

        subgraph "Application Layer"
            API1[API Server 1]
            API2[API Server 2]
            API3[API Server 3]
        end

        subgraph "Cache Layer"
            Redis[(Redis Cache<br/>Session & Data)]
        end

        subgraph "Data Layer"
            MongoDB[(MongoDB<br/>Replica Set)]
            ES[(Elasticsearch<br/>Cluster)]
        end

        subgraph "Event Processing"
            Kafka[Kafka Cluster]
        end
    end
```

```

        Indexer1[Indexer 1]
        Indexer2[Indexer 2]
        Cron[Cron Service]
    end

    Client -->|HTTPS| LB
    LB -->|Round Robin| API1
    LB -->|Round Robin| API2
    LB -->|Round Robin| API3

    API1 -.->|Check Cache| Redis
    API2 -.->|Check Cache| Redis
    API3 -.->|Check Cache| Redis

    API1 -->|Write/Read| MongoDB
    API2 -->|Write/Read| MongoDB
    API3 -->|Write/Read| MongoDB

    API1 -->|Search| ES
    API2 -->|Search| ES
    API3 -->|Search| ES

    API1 -->|Publish| Kafka
    API2 -->|Publish| Kafka
    API3 -->|Publish| Kafka

    Kafka -->|Consume| Indexer1
    Kafka -->|Consume| Indexer2
    Indexer1 -->|Update| ES
    Indexer2 -->|Update| ES

    MongoDB -->|Read| Cron
    Cron -->|Reindex| ES
end

style LB fill:#E67E22
style API1 fill:#4A90E2
style API2 fill:#4A90E2
style API3 fill:#4A90E2
style Redis fill:#DC382D
style MongoDB fill:#47A248
style ES fill:#FEC514
style Kafka fill:#231F20

```

Detailed Request Flow with Cache

```

sequenceDiagram
    participant C as Client
    participant LB as Load Balancer
    participant A as API Server
    participant R as Redis Cache
    participant M as MongoDB
    participant E as Elasticsearch

```

Note over C,E: Get User Profile (Cache Hit)

C->>LB: GET /api/users/123

LB->>A: Forward Request

A->>R: GET user:123

R-->>A: Cache HIT (User Data)

A-->>LB: 200 OK + User Data

LB-->>C: Response

Note over C,E: Get User Profile (Cache Miss)

C->>LB: GET /api/users/456

LB->>A: Forward Request

A->>R: GET user:456

R-->>A: Cache MISS (null)

A->>M: Find user by ID

M-->>A: User Document

A->>R: SET user:456 (TTL: 1h)

R-->>A: OK

A-->>LB: 200 OK + User Data

LB-->>C: Response

Note over C,E: Search Posts (with Cache)

C->>LB: GET /api/posts/search?q=tech

LB->>A: Forward Request

A->>R: GET search:tech:page:1

R-->>A: Cache MISS

A->>E: Search Query

E-->>A: Search Results

A->>R: SET search:tech:page:1 (TTL: 5m)

A-->>LB: 200 OK + Results

LB-->>C: Response

Cache Strategy Details

graph TB

subgraph "Cache Strategy by Data Type"

subgraph "User Data"

U1[User Profile
TTL: 1 hour
Strategy: Cache-Aside]

U2[User Sessions
TTL: 30 min
Strategy: Write-Through]

end

subgraph "Post Data"

P1[Individual Posts
TTL: 30 min
Strategy: Cache-Aside]

P2[Post Lists
TTL: 5 min
Strategy: Cache-Aside]

P3[Trending Posts
TTL: 2 min
Strategy: Refresh-Ahead]

end

subgraph "Search Results"

S1[Search Queries
TTL: 5 min
Strategy: Cache-Aside]

S2[Popular Searches
TTL: 10 min
Strategy: Cache-Aside]

end

subgraph "Computed Data"

C1[Vote Counts
TTL: 1 min
Strategy: Write-Behind]

```

        C2[User Stats<br/>TTL: 15 min<br/>Strategy: Cache-Aside]
    end
end

style U1 fill:#3498DB
style U2 fill:#3498DB
style P1 fill:#2ECC71
style P2 fill:#2ECC71
style P3 fill:#2ECC71
style S1 fill:#F39C12
style S2 fill:#F39C12
style C1 fill:#9B59B6
style C2 fill:#9B59B6

```

Load Balancer Configuration

```

graph LR
    subgraph "Load Balancer Strategies"
        LB[Load Balancer]

        subgraph "Routing Rules"
            R1[/api/users/* → Round Robin]
            R2[/api/posts/* → Round Robin]
            R3[/api/search/* → Least Connections]
            R4[/api/uploads/* → IP Hash]
        end

        subgraph "Health Checks"
            H1[GET /health every 10s]
            H2[Timeout: 5s]
            H3[Unhealthy threshold: 3]
        end

        LB --> R1
        LB --> R2
        LB --> R3
        LB --> R4
        LB -.->|Monitor| H1
        LB -.->|Monitor| H2
        LB -.->|Monitor| H3
    end
end

```

Write Operation with Cache Invalidation

```

sequenceDiagram
    participant C as Client
    participant LB as Load Balancer
    participant A as API Server
    participant R as Redis
    participant M as MongoDB
    participant K as Kafka

    Note over C,K: Update Post Flow

```

```

C->>LB: PUT /api/posts/123
LB->>A: Forward Request

par Database Update
  A->>M: Update Post Document
  M-->>A: Success
and Cache Invalidation
  A->>R: DELETE post:123
  A->>R: DELETE post_list:*
  R-->>A: Keys Deleted
and Event Publishing
  A->>K: Publish "post.updated"
  K-->>A: Acknowledged
end

A-->>LB: 200 OK
LB-->>C: Response

Note over K: Async: Indexer will update Elasticsearch

```

Complete Data Flow

```

flowchart TB
    Start([Client Request]) --> LB[Load Balancer]

    LB -->|Route| API[API Server]

    API --> CheckCache{Check Redis Cache}
    CheckCache -->|HIT| ReturnCache[Return Cached Data]
    CheckCache -->|MISS| CheckType{Request Type}

    CheckType -->|Search| ES[Query Elasticsearch]
    CheckType -->|CRUD| DB[Query MongoDB]

    ES --> CacheResult[Cache Search Results]
    DB --> CacheData[Cache DB Results]

    CacheResult --> Return[Return Response]
    CacheData --> Return
    ReturnCache --> Return

    API -->|Write Op| PublishEvent[Publish to Kafka]
    PublishEvent --> Indexer[Indexer Service]
    Indexer --> UpdateES[Update Elasticsearch]

    API -->|Write Op| InvalidateCache[Invalidate Cache]

    Return --> End([Response to Client])

    style LB fill:#E67E22
    style API fill:#4A90E2
    style ES fill:#FEC514
    style DB fill:#47A248

```

```
style CheckCache fill:#DC382D
style PublishEvent fill:#231F20
```

Comparison: Stage Evolution

Architecture Complexity

```
graph LR
    subgraph "Stage 1"
        S1C[Client] --> S1A[API]
        S1A --> S1D[(DB)]
    end

    subgraph "Stage 2"
        S2C[Client] --> S2A[API]
        S2A --> S2M[(MongoDB)]
        S2A --> S2E[(Elasticsearch)]
        S2A --> S2K[Kafka]
        S2K --> S2I[Indexer]
        S2I --> S2E
    end

    subgraph "Stage 3"
        S3C[Client] --> S3L[LB]
        S3L --> S3A1[API-1]
        S3L --> S3A2[API-2]
        S3A1 --> S3R[(Redis)]
        S3A2 --> S3R
        S3A1 --> S3M[(MongoDB)]
        S3A2 --> S3M
        S3A1 --> S3E[(ES)]
        S3A2 --> S3E
    end

    style S1A fill:#4A90E2
    style S2A fill:#4A90E2
    style S3A1 fill:#4A90E2
    style S3A2 fill:#4A90E2
    style S3L fill:#E67E22
    style S3R fill:#DC382D
```

Performance Metrics Comparison

```
graph TB
    subgraph "Metrics Evolution"
        direction TB
        M1["Stage 1<br/>---<br/>Response Time: 200ms<br/>Throughput: 100 req/s<br/>Availability: 95%"]
        M2["Stage 2<br/>---<br/>Response Time: 150ms<br/>Search Time: 50ms<br/>Throughput: 300 req/s<br/>"]
    end
```



```

M3[Stage 3<br/>---<br/>Response Time: 50ms avg<br/>Search Time: 30ms<br/>Throughput: 2000 req/s

M1 -.->|Added Search| M2
M2 -.->|Added LB & Cache| M3
end

style M1 fill:#E74C3C
style M2 fill:#F39C12
style M3 fill:#2ECC71

```

Key Features by Stage

Feature	Stage 1	Stage 2	Stage 3
User CRUD			
Post CRUD			
Full-Text Search			
Search Ranking		(Relevance, Recency, Votes)	
Real-time Indexing		(Kafka)	
Batch Reindexing		(Cron)	
Load Balancing			
Caching			(Redis)
Horizontal Scaling		Partial	
High Availability			

Technology Stack Summary

Technology Stack			
Layer	Stage 1	Stage 2	Stage 3
Load Bal.	-	-	Nginx/ALB
API	Node.js	Node.js	Node.js (3x)
Cache	-	-	Redis
Database	PostgreSQL	MongoDB	MongoDB Cluster
Search	-	Elasticsearch	ES Cluster
Queue	-	Kafka	Kafka Cluster
Indexer	-	Node.js	Node.js (2x)
Scheduler	-	Cron	Cron Service